



Bridge of Life
Education

FINN Compiler

Lecturer: Hua-Yang Weng

Date: 2022/08/18

[\[FINN-R: An End-to-End Deep-Learning Framework for Fast Exploration of Quantized Neural Networks\]](#)

[FPGA'17: FINN: A Framework for Fast, Scalable Binarized Neural Network Inference]
(<https://arxiv.org/abs/1612.07119>)

From AI to Gate Textbook

From AI to Gate

Preface

Getting Started

Network Define

Compiler

Introduction

Tidy-up and Preprocess

Streamlining

HLS Dataflow

Hardware Build

Verification

NN Hardware

HLS

Case Study

Code Repository

Translations

Compiler

In this chapter, we are going to explain how FINN maps the deep learning network from an exported ONNX graph into high-level-synthesized layers. From the [Wikipedia definition](#), the name "compiler" is primarily used for programs that translate source code from a high-level programming language to a lower level language (e.g. assembly language, object code, or machine code) to create an executable program. However, what the FINN Compiler here actually transforms is **from high-level operation representations such as ONNX graph, into another operation representation that is compatible with some high-level-synthesis libraries**. The later will perform C/C++ function-calls to the HLS library and then be synthesized by another High-Level-Synthesis compiler such as Vitis-HLS (into hardware representations and then the bitstream file for FPGAs).

Goals

The figure below shows our objective in this chapter. As chapter one has shown, we are going to transform the onnx graph from the left to the right. We can see that all the operations are now mapped into hardware operations supported by certain HLS hardware unit.

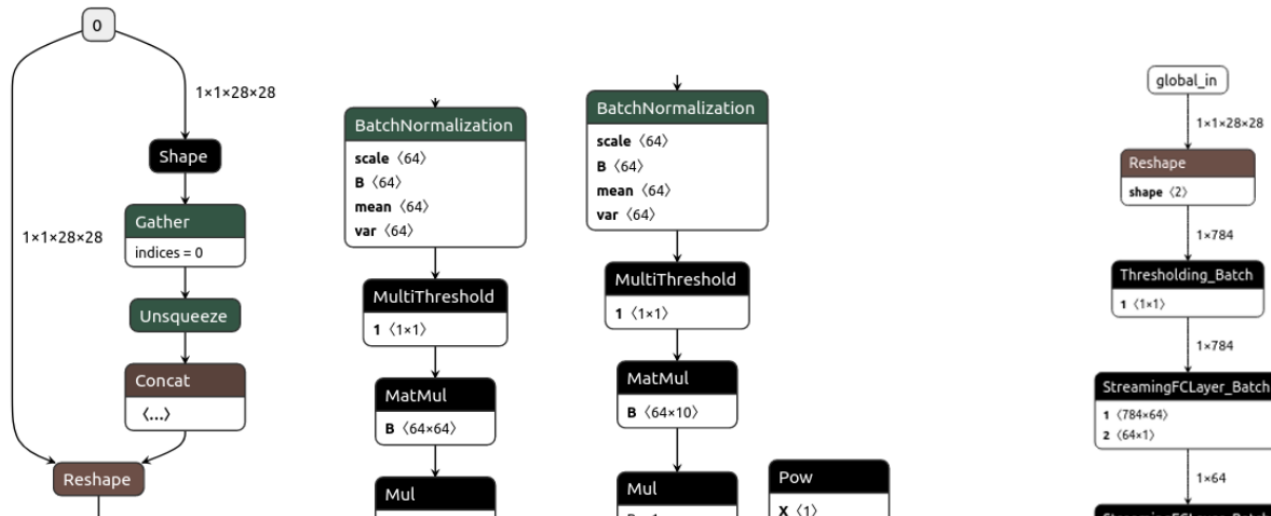


Table of Contents

Compiler

Goals

Chapter structure

Overview

- End-to-End Compiling
- Phase (I): Brevitas export
- Phase (II):
 - Network preparation
 - Conversion to HLS layers
- Phase (III): Hardware Build
- Phase (IV): PYNQ deployment

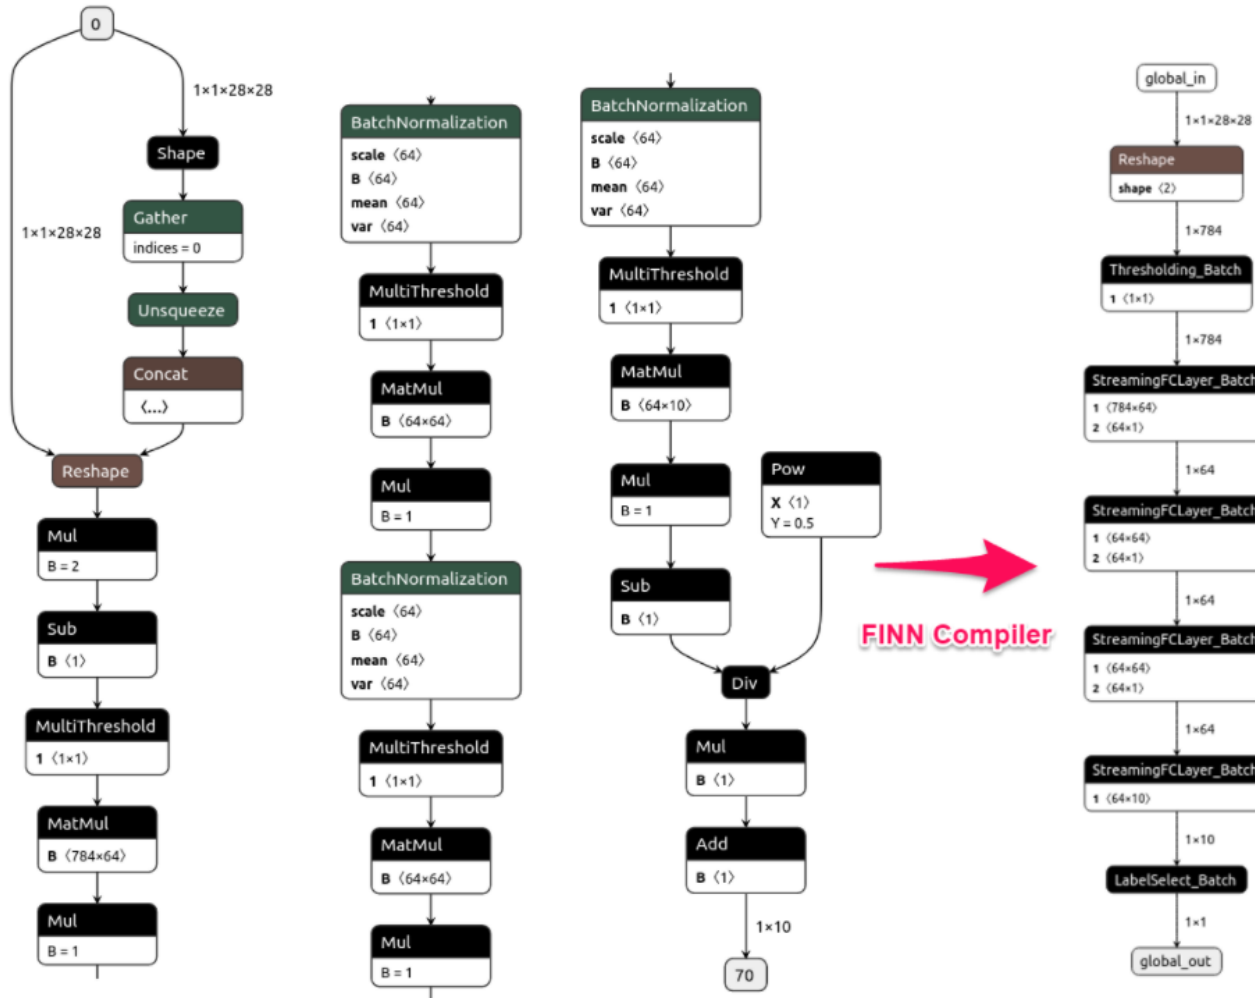
Overview

- End-to-End Compiling
- Phase (I): Brevitas export
- Phase (II):
 - Network preparation
 - Conversion to HLS layers
- Phase (III): Hardware Build
- Phase (IV): PYNQ deployment

Introduction

- The FINN compiler comes with *many transformations* that modify the ONNX representation of the network according to certain patterns.
- Phase (I): Brevitas export
- Phase (II):
 - Network preparation
 - Conversion to HLS layers
- Phase (III): Hardware Build
- Phase (IV): PYNQ deployment

Goal



Overview

- End-to-End Compiling
- **Phase (I): Brevitas export**
- Phase (II):
 - Network preparation
 - Conversion to HLS layers
- Phase (III): Hardware Build
- Phase (IV): PYNQ deployment



Phase (I): Brevitas Export

Hardware support only for layers defined in Finn-HLS & **Compiler**

Phase(I)

Brevitas Export

Trained Network in Pytorch/Brevitas

Brevitas FINN-ONNX Export

Network Preparation

Network of high-level ONNX layers

Streamlining Transformations

Streamlined network of high-level ONNX layers

Convert to HLS Layers

Mixture of HLS and non-HLS layers

Dataflow Partitioning

Parent Graph (non-HLS layers)

Network of HLS layers, maximum folding

Adjust folding to maximize performance

Network of HLS layers, desired folding

Hardware Build

Create HLS IP per layer

Network of HLS layers, IP per layer

Create stitched design

Network of HLS layers, stitched IP

Create Vivado/Vitis project

Network of HLS layers + Vivado/Vitis project

Synthesis, P & R

Generate PYNQ runtime (driver)

Bitfile

PYNQ Python driver

Run on Hardware

Simulation & Emulation Flows

Simulation using Python

Prepare cppsim

Network of HLS layers with C++ wrappers

Simulation (cppsim) using C++

Prepare rtlsim (stitched)

Prepare rtlsim (node-by-node)

Full network Verilog model

Network of HLS layers with Verilog models

Emulation (rtlsim) using PyVerilog

Don't touch:
Since we are not concern of onnx operations

Coding (to-hls):
Infer_XXX.py
(adjusting nodes)

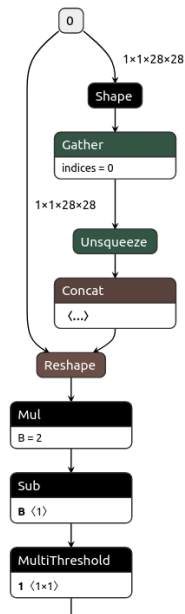
Coding: (HLS Custom Op):
xxx.py
(gen_code)

Coding: (HLS Library):
xxx.cpp/.h (core)

Phase (I): Brevitas Export (1/2)

```
import onnx
from finn.util.test import get_test_model_trained
import brevitas.onnx as bo

tfc = get_test_model_trained("TFC", 1, 1)
bo.export_finn_onnx(tfc, (1, 1, 28, 28), build_dir+"/tfc_w1_a1.onnx")
```



MODEL PROPERTIES	
format	ONNX v6
producer	pytorch 1.9
imports	ai.onnx v9
INPUTS	
0	name: 0 type: float32[1, 1, 28, 28]
OUTPUTS	
70	name: 70 type: float32[1, 10]

Phase (I): Brevitas Export (2/2)

- **ModelWrapper**

- Wrapper around the ONNX model which provides several helper functions to make it easier to work with the model.

```
from finn.core.modelwrapper import ModelWrapper  
model = ModelWrapper(build_dir+"/tfc_w1_a1.onnx")
```

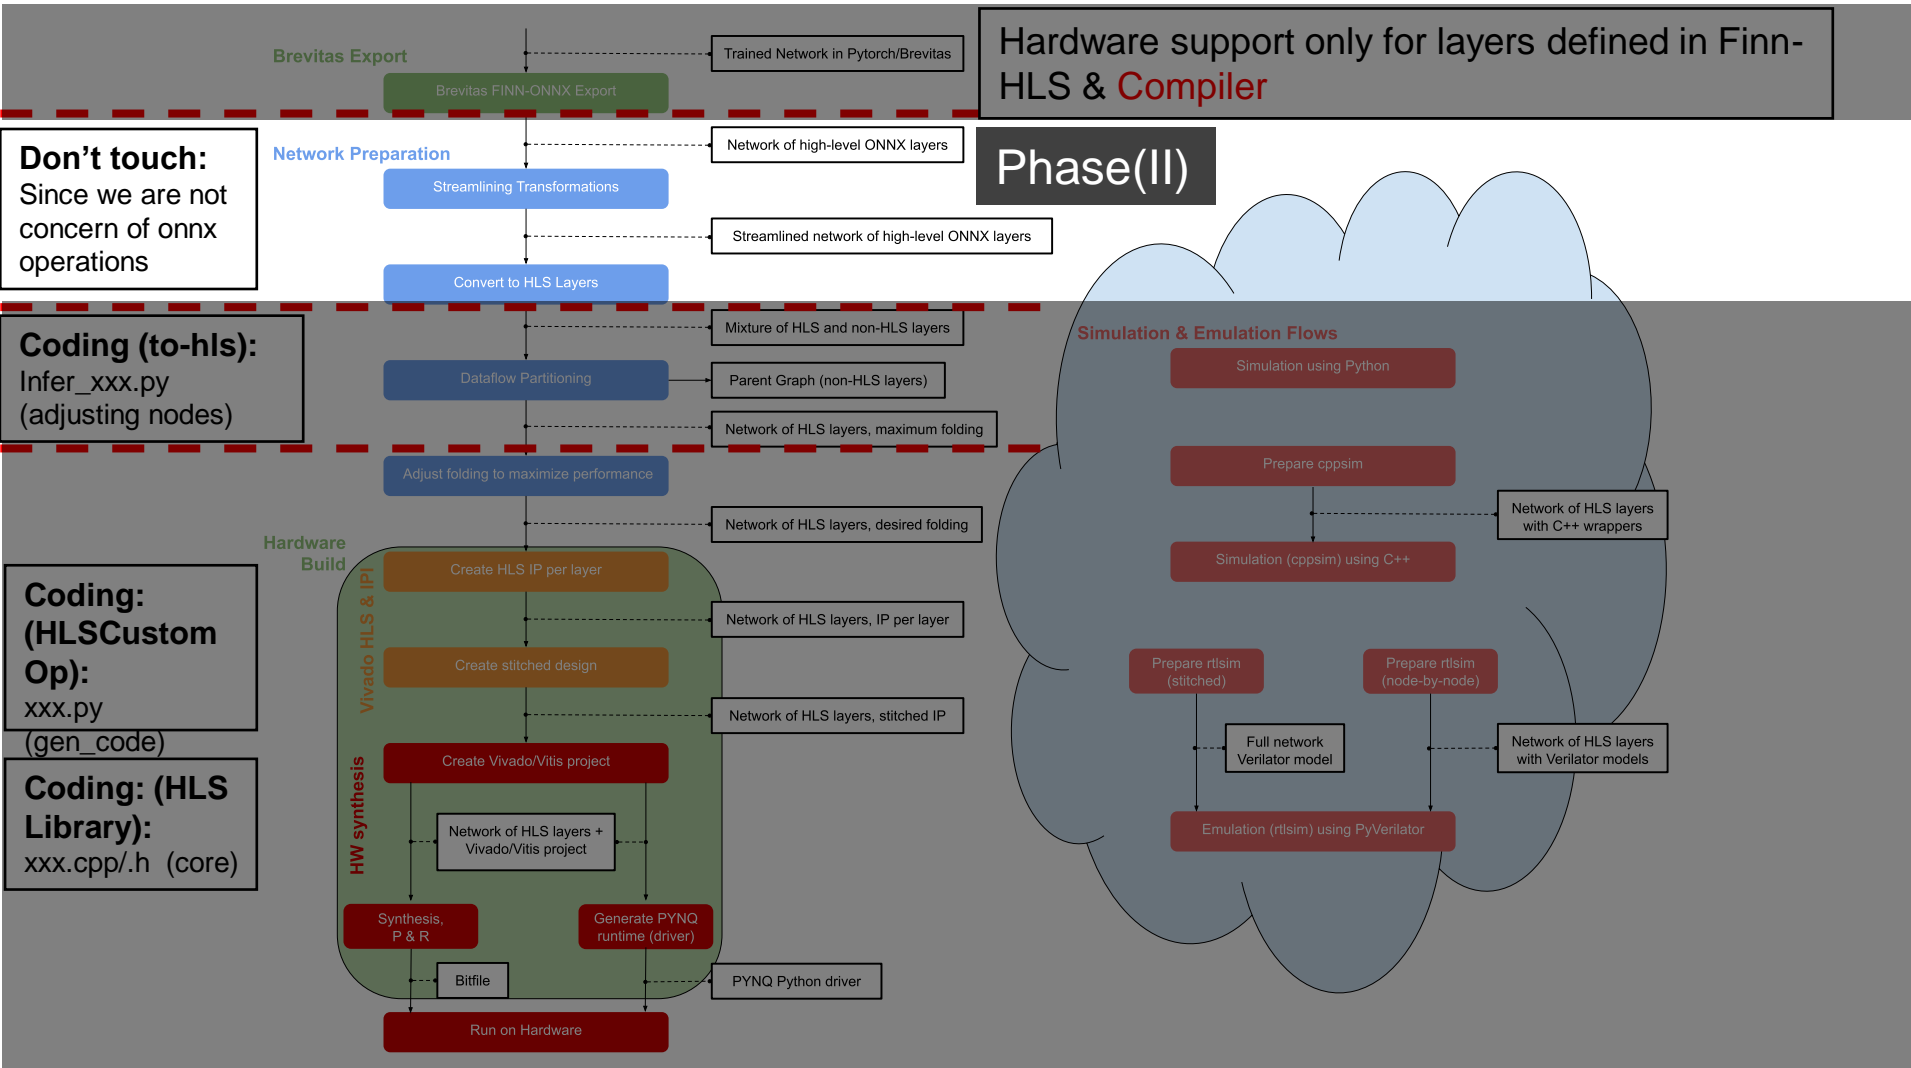
- Principle of FINN: **Analysis and Transformation passes**

- **Analysis pass:** extracts specific information about the model.
- **Transformation pass:** changes the model and returns the changed model

Overview

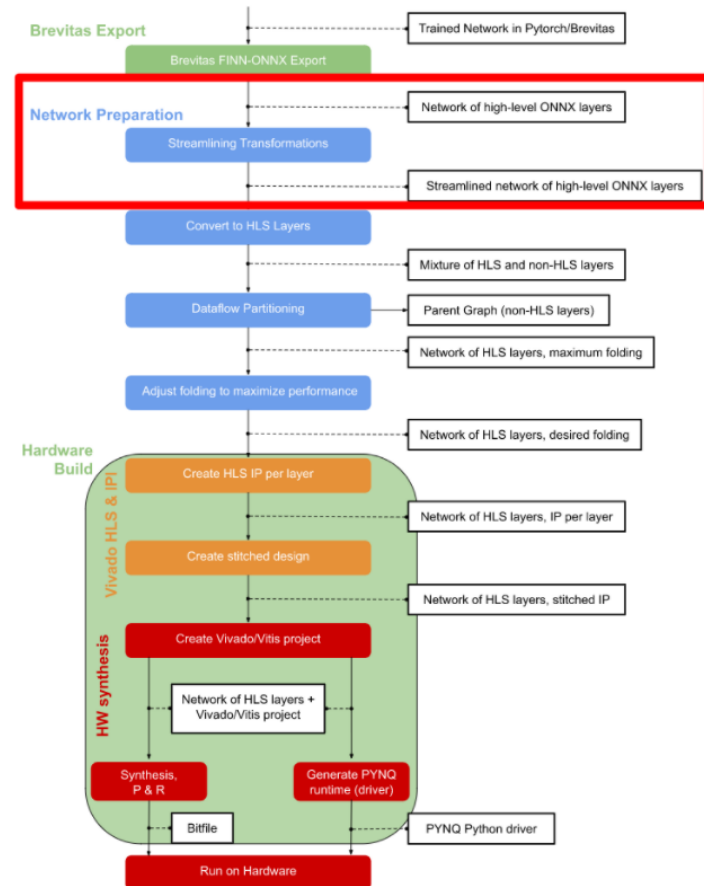
- End-to-End Compiling
- Phase (I): Brevitas export
- Phase (II):
 - Network preparation
 - Conversion to HLS layers
- Phase (III): Hardware Build
- Phase (IV): PYNQ deployment

Phase (II): Network preparation



Phase (II): Network preparation

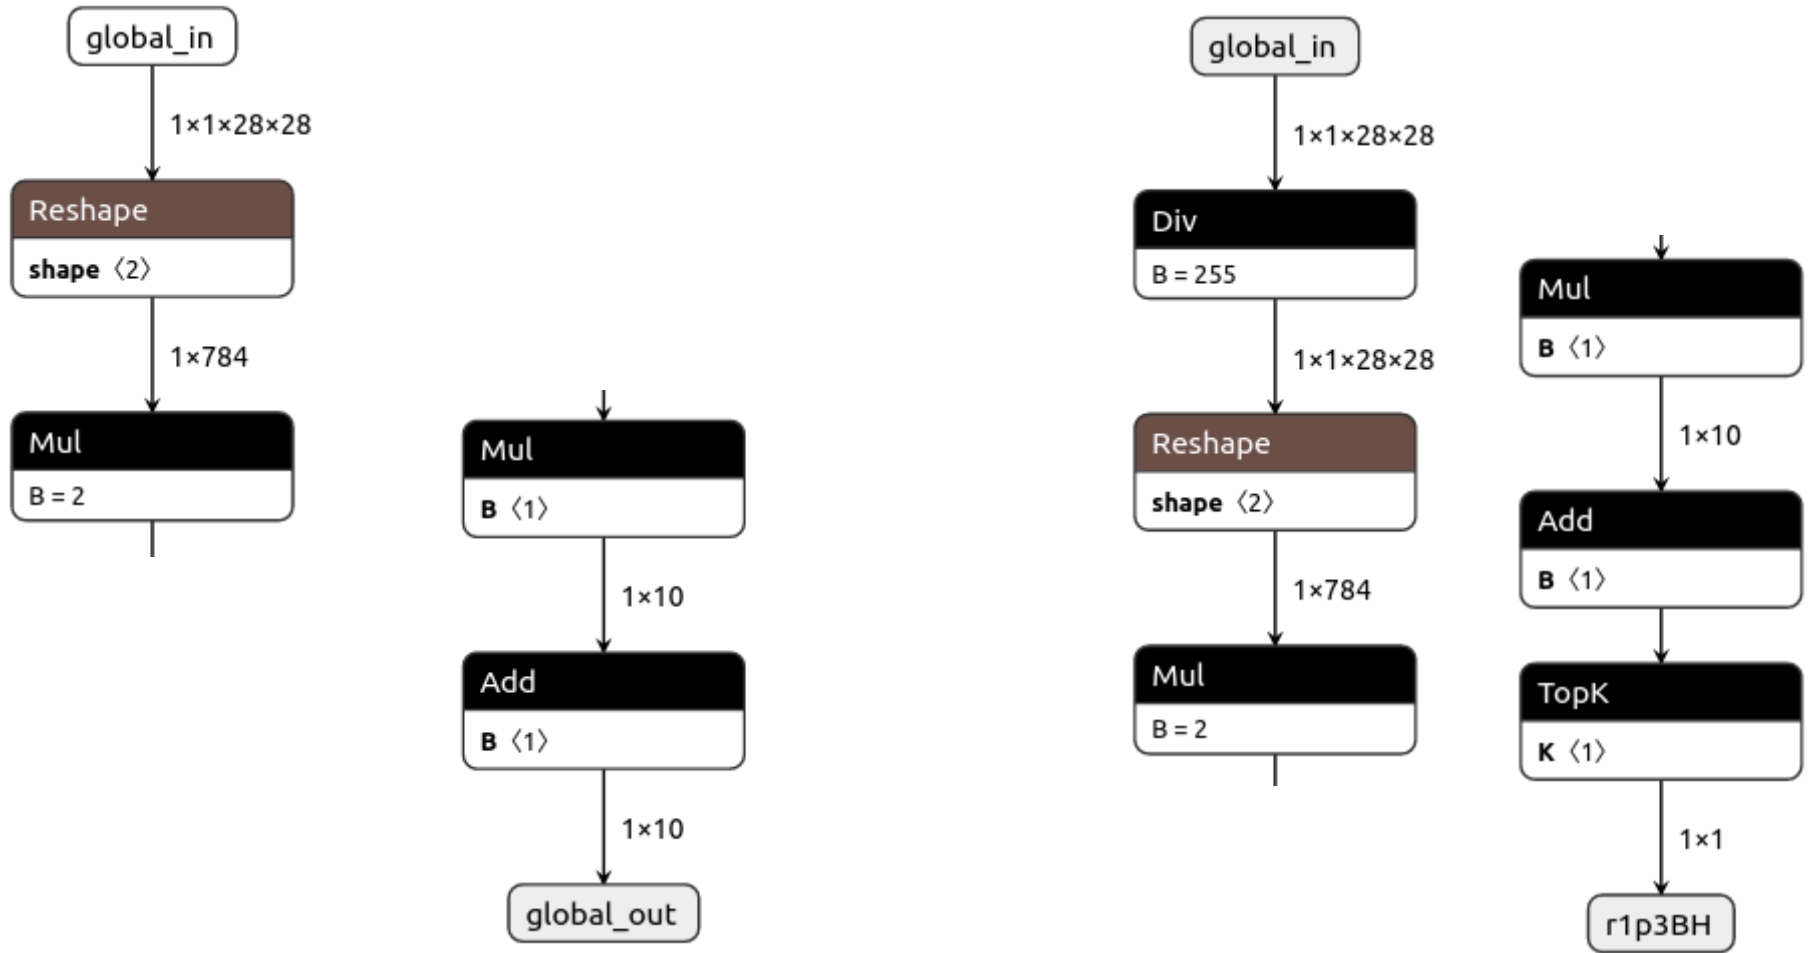
- (I) FINN-style Dataflow Architectures
- (II) Tidy-up transformations
- (III) Adding Pre- and Postprocessing
- (IV) Streamlining
- (V) Conversion to HLS layers
- (VI) Creating a Dataflow Partition
- (VII) Folding: Adjusting the Parallelism



(II) Tidy-up transformations

- GiveUniqueNodeNames
- GiveReadableTensorNames
- InferShapes
- InferDataTypes
- FoldConstants
- RemoveStaticGraphInputs

(III) Adding Pre- and Postprocessing



(IV) Streamlining(1/4)

- **Goal:** eliminate, collapsing floating point operations

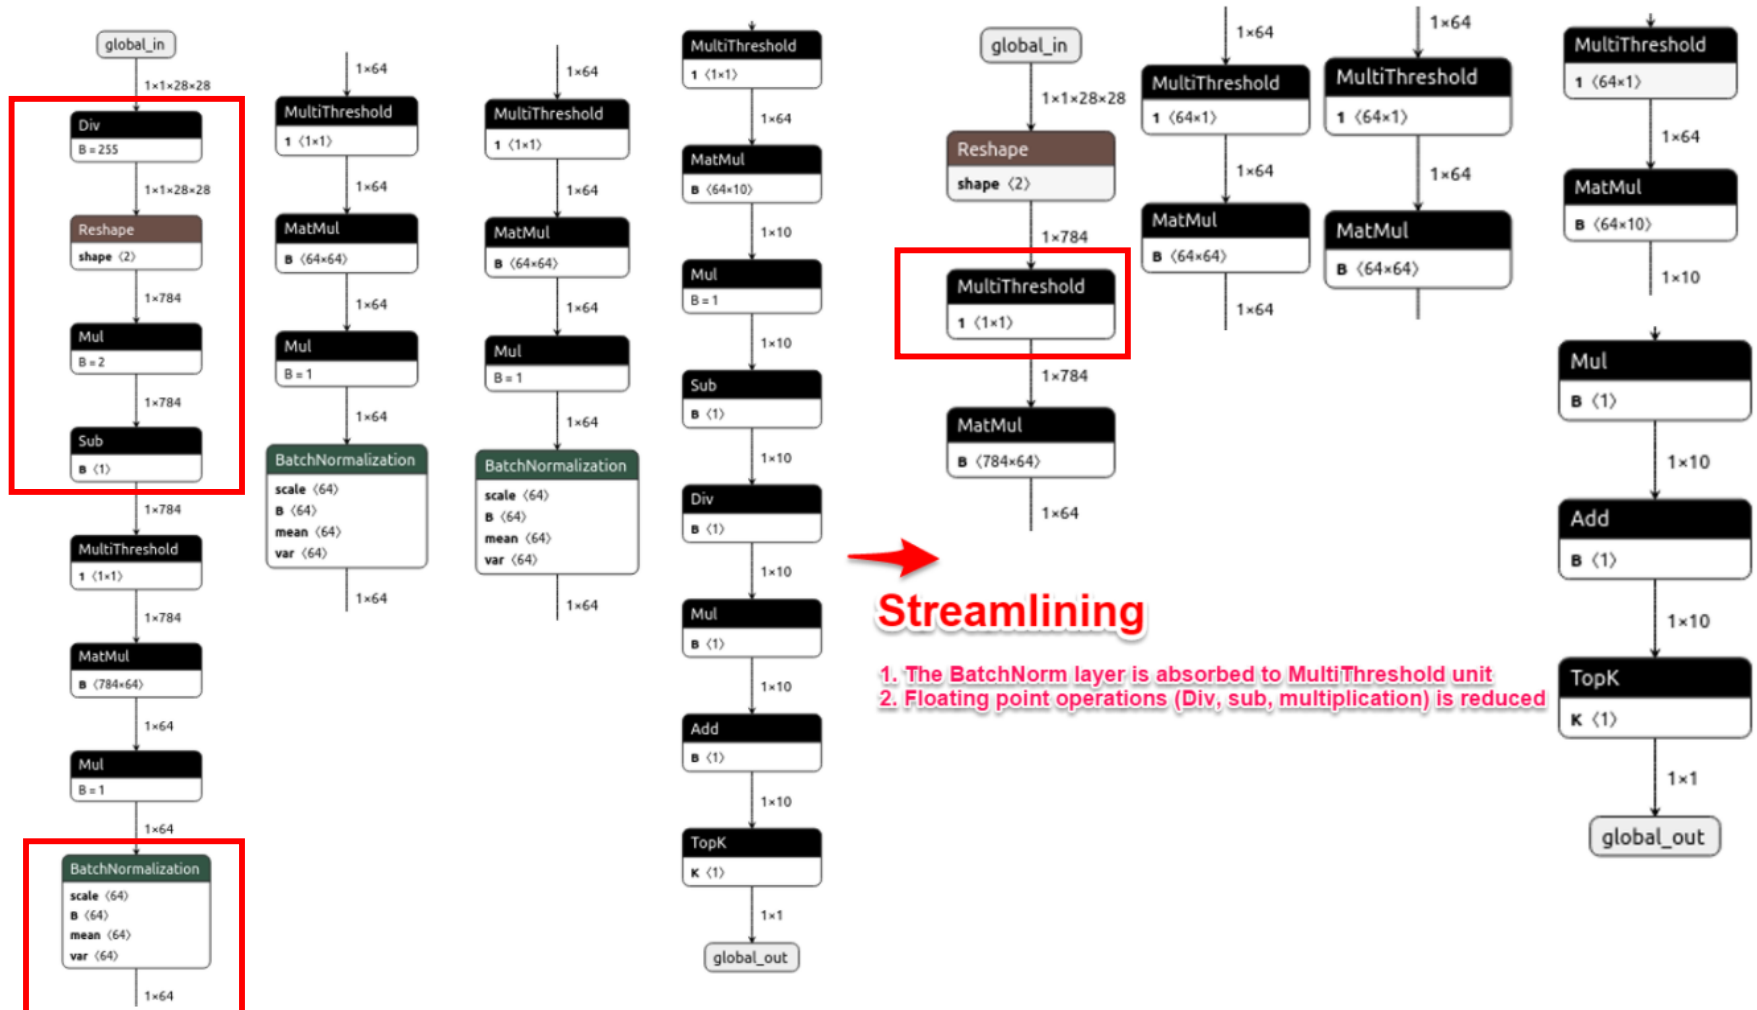
```
from finn.transformation.streamline.reorder import MoveScalarLinearPastInvariants
import finn.transformation.streamline.absorb as absorb

model = ModelWrapper(build_dir+"/tfc_w1_a1_pre_post.onnx")
# move initial Mul (from preproc) past the Reshape
model = model.transform(MoveScalarLinearPastInvariants())
# streamline
model = model.transform(Streamline())
model.save(build_dir+"/tfc_w1_a1_streamlined.onnx")
showInNetron(build_dir+"/tfc_w1_a1_streamlined.onnx")
```

```
class Streamline(Transformation):
    """Apply the streamlining transform, see arXiv:1709.04060."""

    def apply(self, model):
        streamline_transformations = [
            ConvertSubToAdd(),
            ConvertDivToMul(),
            BatchNormToAffine(),
            ConvertSignToThres(),
            AbsorbSignBiasIntoMultiThreshold(),
            MoveAddPastMul(),
            MoveScalarAddPastMatMul(),
            MoveAddPastConv(),
            MoveScalarMulPastMatMul(),
            MoveScalarMulPastConv(),
            MoveAddPastMul(),
            CollapseRepeatedAdd(),
            CollapseRepeatedMul(),
            AbsorbAddIntoMultiThreshold(),
            FactorOutMulSignMagnitude(),
            AbsorbMulIntoMultiThreshold(),
            Absorb1BitMulIntoMatMul(),
            Absorb1BitMulIntoConv(),
            RoundAndClipThresholds(),
        ]
        for trn in streamline_transformations:
            model = model.transform(trn)
        model = model.transform(RemoveIdentityOps())
        model = model.transform(GiveUniqueNodeNames())
        model = model.transform(GiveReadableTensorNames())
        model = model.transform(InferDataTypes())
        return (model, False)
```


(IV) Streamlining(2/4)



(IV) Streamlining(3/4)

- Current implementation of streamlining:
 - Highly network-specific (topology change)

```
from finn.transformation.bipolar_to_xnor import ConvertBipolarMatMulToXnorPopcount
from finn.transformation.streamline.round_thresholds import RoundAndClipThresholds
from finn.transformation.infer_data_layouts import InferDataLayouts
from finn.transformation.general import RemoveUnusedTensors

model = model.transform(ConvertBipolarMatMulToXnorPopcount())
model = model.transform(absorb.AbsorbAddIntoMultiThreshold())
model = model.transform(absorb.AbsorbMulIntoMultiThreshold())
# absorb final add-mul nodes into TopK
model = model.transform(absorb.AbsorbScalarMulAddIntoTopK())
model = model.transform(RoundAndClipThresholds())

# bit of tidy-up
model = model.transform(InferDataLayouts())
model = model.transform(RemoveUnusedTensors())

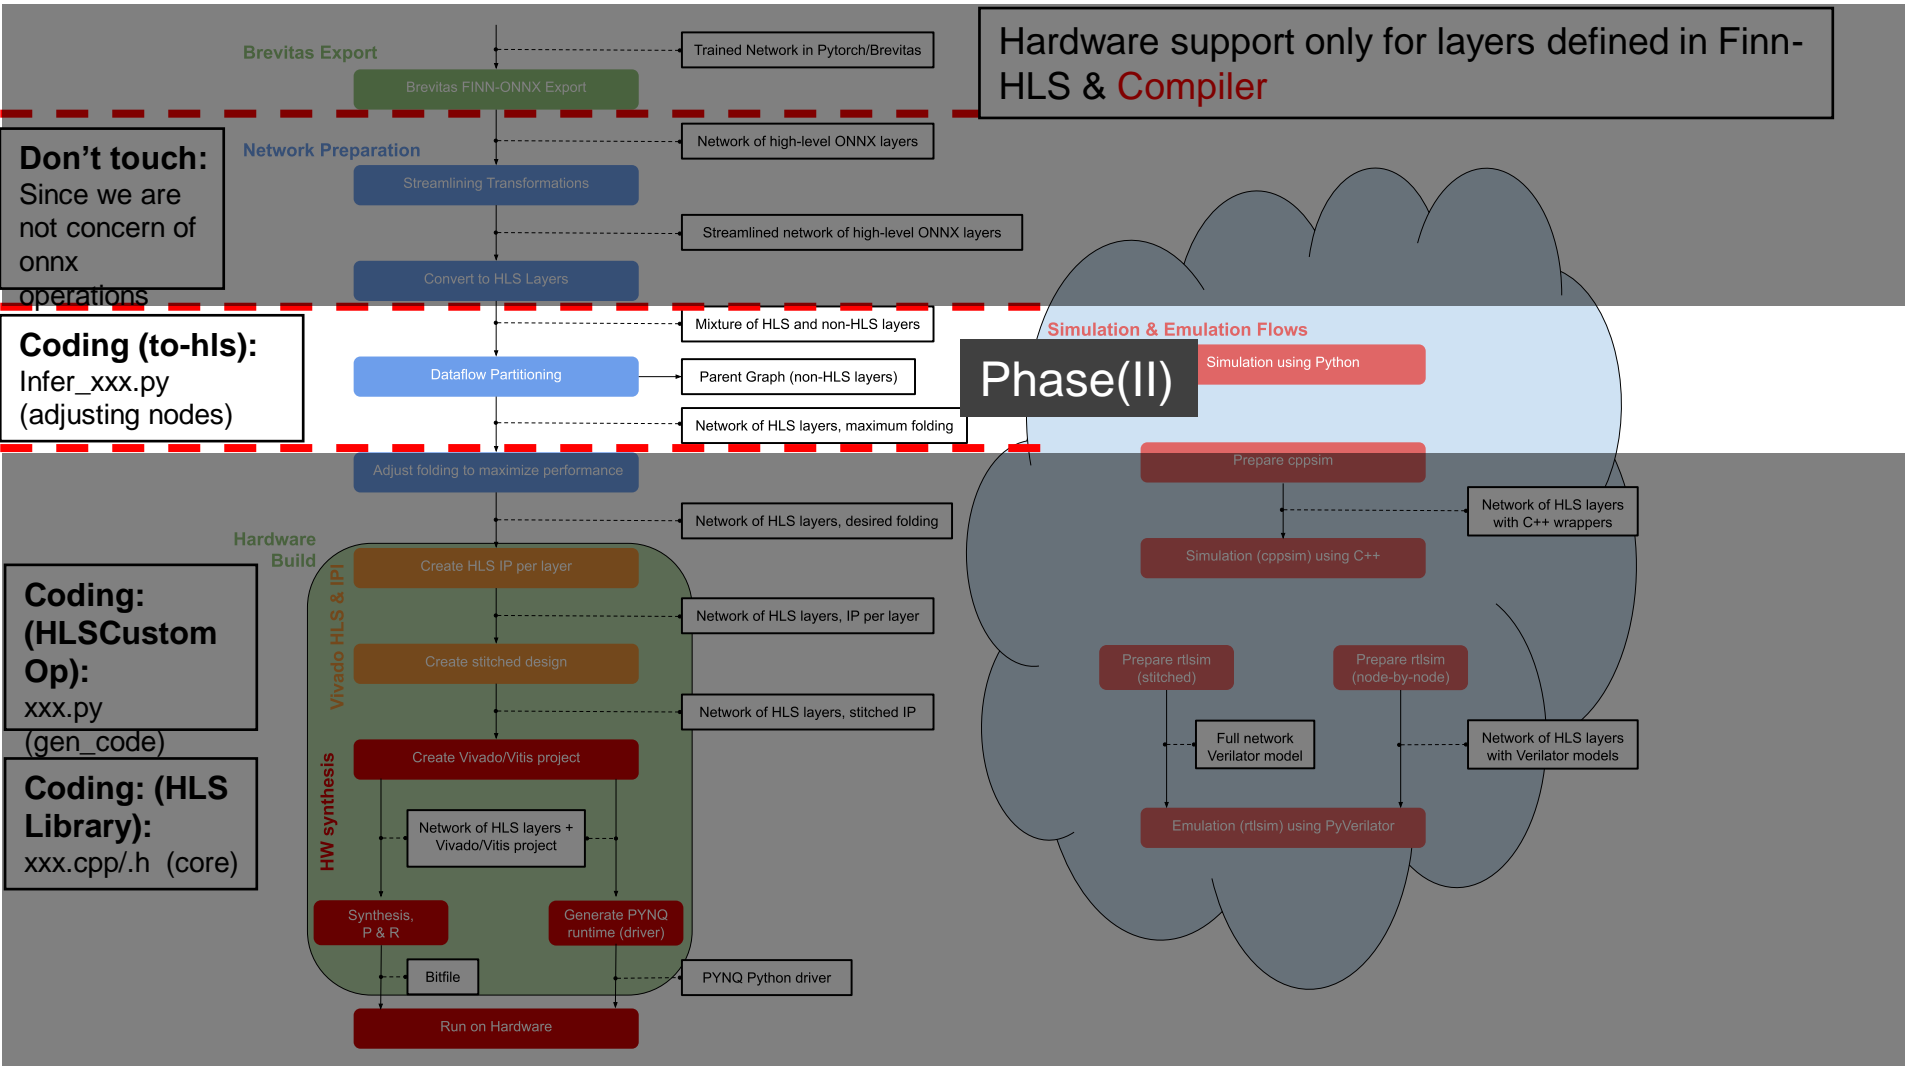
model.save(build_dir+"/tfc_wla1_ready_for_hls_conversion.onnx")
showInNetron(build_dir+"/tfc_wla1_ready_for_hls_conversion.onnx")
```



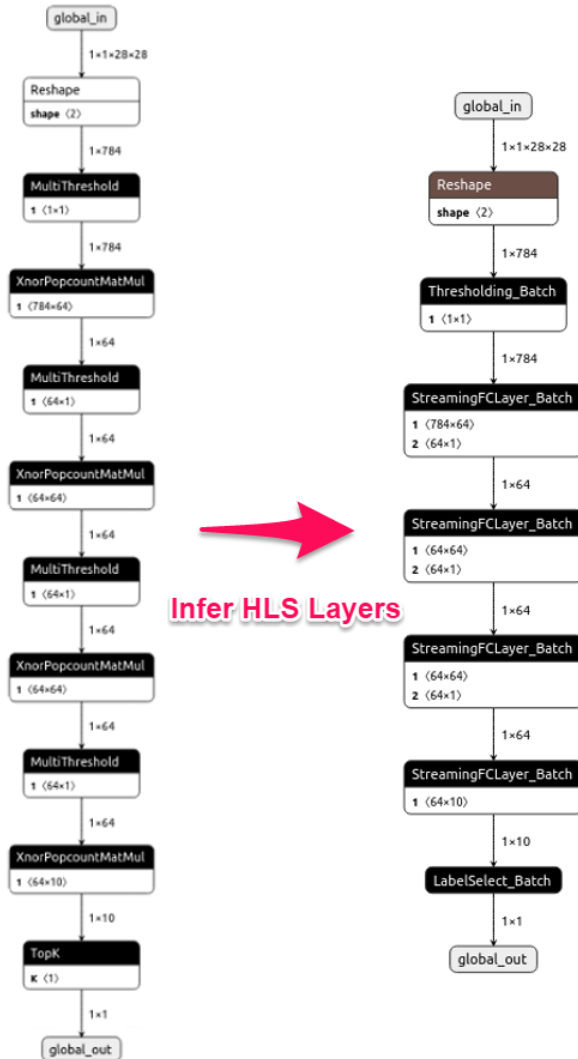
Overview

- End-to-End Compiling
- Phase (I): Brevitas export
- **Phase (II):**
 - Network preparation
 - Conversion to HLS layers
- Phase (III): Hardware Build
- Phase (IV): PYNQ deployment

Phase (II): (V) Conversion to HLS layers



(V) Conversion to HLS layers



Infer HLS Layers

NODE PROPERTIES

type StreamingFCLayer_Batch
domain finn.custom_op.fpgadataflow



ATTRIBUTES

accDataType UINT10
ActVal 0
backend fpgadataflow
binaryXnorMode 1
inputDataType BINARY
mem_mode decoupled
MH 64
MW 784
noActivation 0
numInputVectors 1
outputDataType BINARY
PE 1
SIMD 1
weightDataType BINARY


INPUTS

0 name: MultiThreshold_0_out0
1 name: MatMul_0_param0
2 name: MultiThreshold_1_param0

OUTPUTS



■ HLS components
■ Verilog components



NODE PROPERTIES

type StreamingFCLayer_Batch
domain finn.custom_op.fpgadataflow

ATTRIBUTES

accDataType UINT10

ActVal 0

backend fpgadataflow

binaryXnorMode 1

inputDataType BINARY

mem_mode decoupled

MH 64

MW 784

noActivation 0

numInputVectors 1

outputDataType BINARY

PE 1


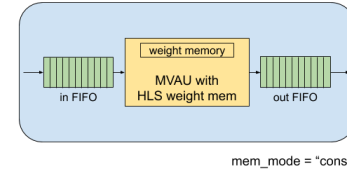
SIMD 1

weightDataType BINARY

INPUTS

0 name: MultiThreshold_0_out0
1 name: MatMul_0_param0
2 name: MultiThreshold_1_param0

OUTPUTS



Legend:
HLS components (yellow)
Verilog components (green)

mem_mode:

- const: weights are “baked in” into the HLS code
- decoupled: weights are streamed into the core using streamers and fifos

(See FINN doc for details:
<https://finn.readthedocs.io/en/latest/internals.html>)

(V) Conversion to HLS layers: Infer_xxx.py

```
import finn.transformation.fpgadataflow.convert_to_hls_layers as to_hls
model = ModelWrapper(build_dir+"/tfc_w1a1_ready_for_hls_conversion.onnx")
model = model.transform(to_hls.InferBinaryStreamingFCLayer("decoupled"))
# TopK to LabelSelect
model = model.transform(to_hls.InferLabelSelectLayer())
# input quantization (if any) to standalone thresholding
model = model.transform(to_hls.InferThresholdingLayer())
model.save(build_dir+"/tfc_w1_a1_hls_layers.onnx")
showInNetron(build_dir+"/tfc_w1_a1_hls_layers.onnx")
```

```
class InferBinaryStreamingFCLayer(Transformation):
```

```
    """Convert XnorPopcountMatMul layers to
    StreamingFCLayer_Batch layers. Any immediately following MultiThreshold
    layers will also be absorbed into the MVTU."""
```

```
    def __init__(self, mem_mode="const"):
        super().__init__()
        self.mem_mode = mem_mode
```

```
    def apply(self, model):
```

```
        graph = model.graph
        node_ind = 0
        graph_modified = False
        for n in graph.nodes:
```

```
            node_ind += 1
            if n.op_type == "XnorPopcountMatMul":
```

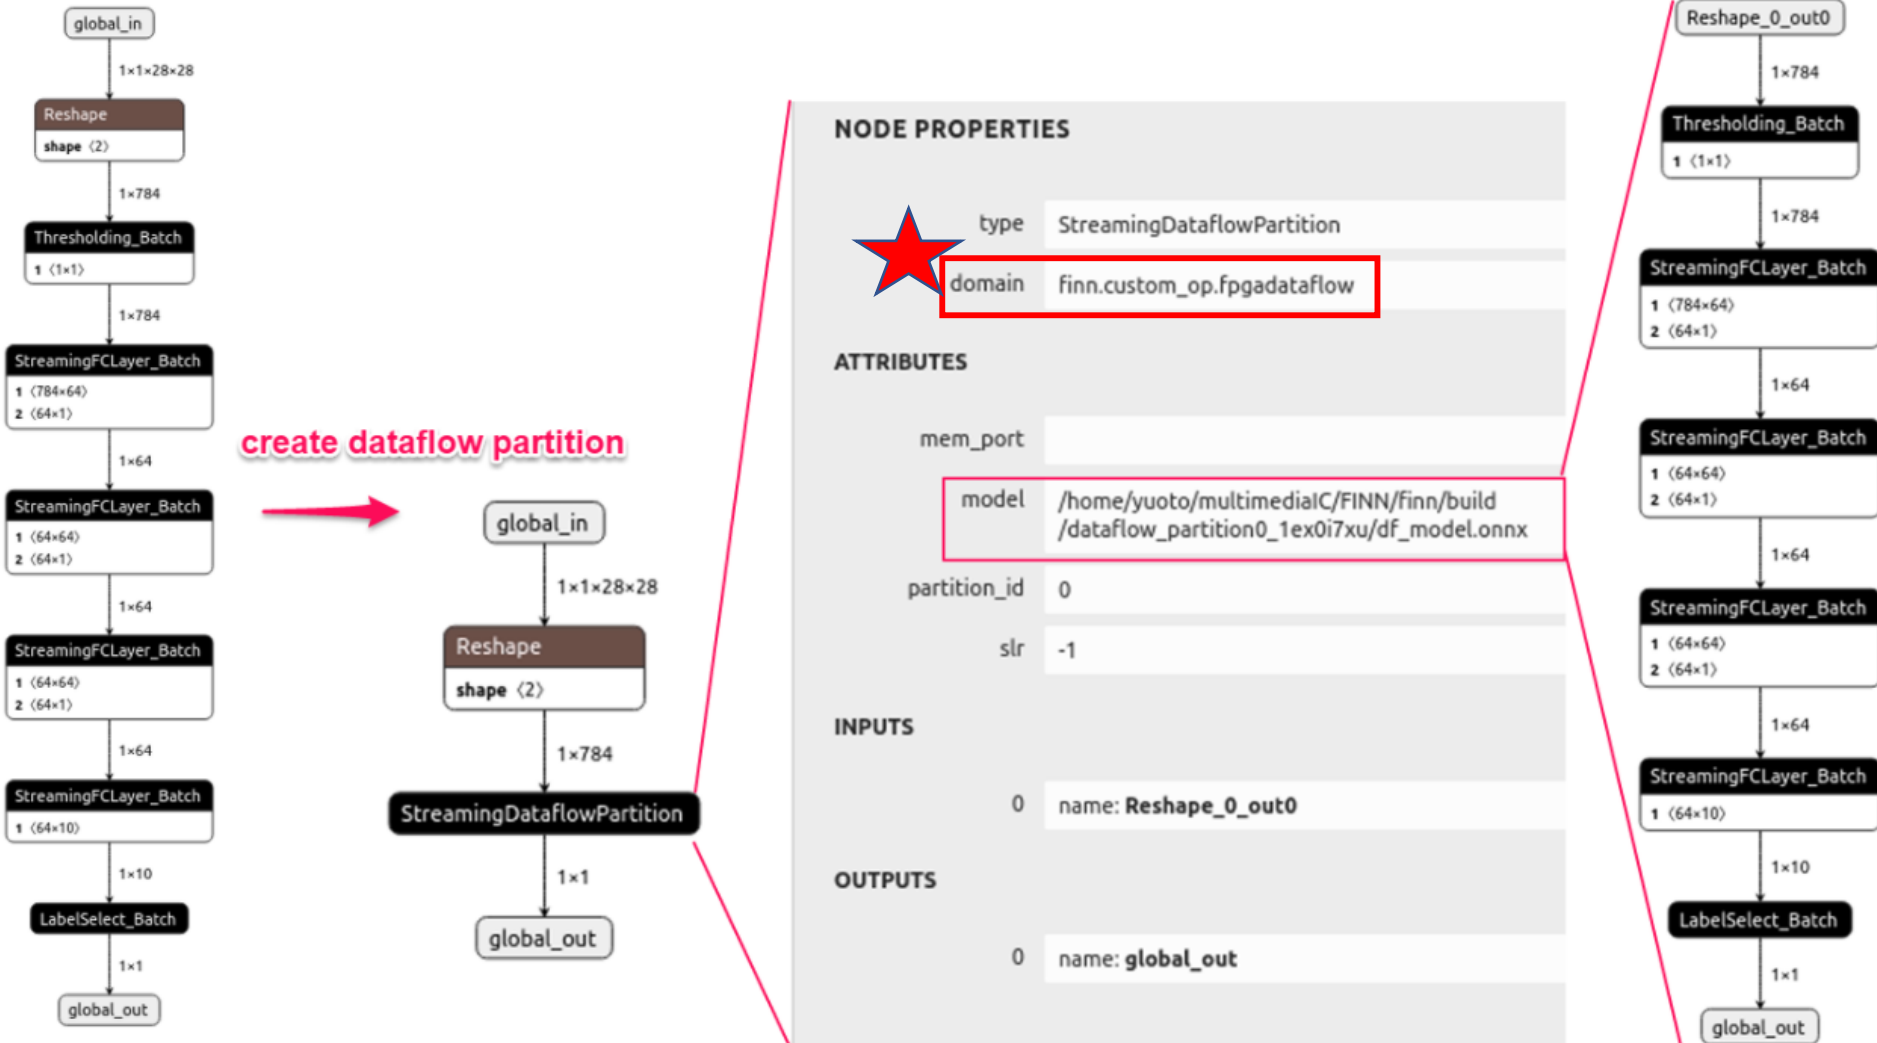
```
                mm_input = n.input[0]
                mm_weight = n.input[1]
                mm_output = n.output[0]
                mm_in_shape = model.get_tensor_shape(mm_input)
                mm_out_shape = model.get_tensor_shape(mm_output)
                assert (
                    model.get_tensor_datatype(mm_input) == DataType.BINARY
                ), "First
```

```
    # Create and insert new StreamingFCLayer node
```

```
    new_node = helper.make_node(
        "StreamingFCLayer_Batch",
        [mm_input, mm_weight, mt_thres],
        [mt_output],
        domain="finn.custom_op.fpgadataflow",
        backend="fpgadataflow",
        MW=mm_weight,
        MH=mt_thres,
        SIMD=simd,
        PE=pe,
        inputDataTypes=[mm_input.name, mm_weight.name],
        outputDataTypes=[mt_output.name],
        ActVal=actval,
        binaryXnorMode=1,
        noActivation=0,
        numInputVectors=list(mm_in_shape[:-1]),
        mem_mode=self.mem_mode,
    )
```

```
    graph.node.insert(node_ind, new_node)
    # remove old nodes
```


(VI) Creating a Dataflow Partition



(VII) Folding: Adjusting the Parallelism (1/3)

- Manually set the folding factors and FIFO depths
- Automatically tune parameters:
 - Given an FPGA resource budget
 - Analytical model from the FINN-R paper.

CustomOp wrapper is of class Thresholding_Batch

```
{'PE': ('i', True, 0),
'NumChannels': ('i', True, 0),
'ram_style': ('s', False, 'distributed'),
'inputDataType': ('s', True, ''),
'outputDataType': ('s', True, ''),
'inFIFOdepth': ('i', False, 2),
'outFIFOdepth': ('i', False, 2),
'numInputVectors': ('ints', False, [1]),
'ActVal': ('i', False, 0),
'backend': ('s', True, 'fpgadataflow'),
'code_gen_dir_cppsim': ('s', False, ''),
'code_gen_dir_ipgen': ('s', False, ''),
'executable_path': ('s', False, ''),
'ipgen_path': ('s', False, ''),
'ip_path': ('s', False, ''),
'ip_vlnv': ('s', False, ''),
'exec_mode': ('s', False, ''),
'cycles_rtlsim': ('i', False, 0),
'cycles_estimate': ('i', False, 0),
'rtlsim_trace': ('s', False, ''),
'res_estimate': ('s', False, ''),
'res_hls': ('s', False, ''),
'res_synth': ('s', False, ''),
'rtlsim_so': ('s', False, ''),
'partition_id': ('i', False, 0)}
```

(VII) Folding: Adjusting the Parallelism (2/3)

- Below:

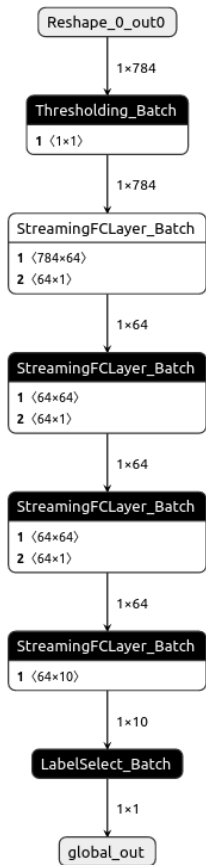
Will be added later with 'InsertFIFO()' Transformation

Set the PE and SIMD s.t. $II = 64$ for each layer

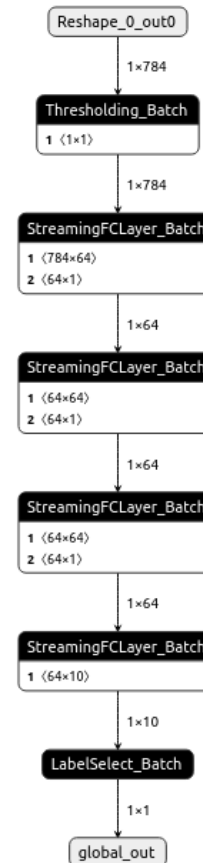
```
fc_layers = model.get_nodes_by_op_type("StreamingFCLayer_Batch")
# (PE, SIMD, in_fifo_depth, out_fifo_depth, ramstyle) for each layer
config = [
    (16, 49, 16, 64, "block"),
    (8, 8, 64, 64, "auto"),
    (8, 8, 64, 64, "auto"),
    (10, 8, 64, 10, "distributed"),
]
for fcl, (pe, simd, ififo, ofifo, ramstyle) in zip(fc_layers, config):
    fcl_inst = getCustomOp(fcl)
    fcl_inst.set_nodeattr("PE", pe)
    fcl_inst.set_nodeattr("SIMD", simd)
    fcl_inst.set_nodeattr("inFIFOdepth", ififo)
    fcl_inst.set_nodeattr("outFIFOdepth", ofifo)
    fcl_inst.set_nodeattr("ram_style", ramstyle)

# set parallelism for input quantizer to be same as first layer's SIMD
inp_qnt_node = model.get_nodes_by_op_type("Thresholding_Batch")[0]
inp_qnt = getCustomOp(inp_qnt_node)
inp_qnt.set_nodeattr("PE", 49)
```

(VII) Folding: Adjusting the Parallelism (3/3)



NODE PROPERTIES	
type	StreamingFCLayer_Batch
module	finn.custom_op.fpgadataflow
ATTRIBUTES	
accDataType	UINT10
ActVal	0
backend	fpgadataflow
binaryXnorMode	1
inputDataType	BINARY
mem_mode	decoupled
MH	64
MW	784
noActivation	0
numInputVectors	1
outputDataType	BINARY
PE	1
SIMD	1
weightDataType	BINARY
INPUTS	
0	name: MultiThreshold_0_out0
1	name: MatMul_0_param0
2	name: MultiThreshold_1_param0
OUTPUTS	
0	name: MultiThreshold_1_out0

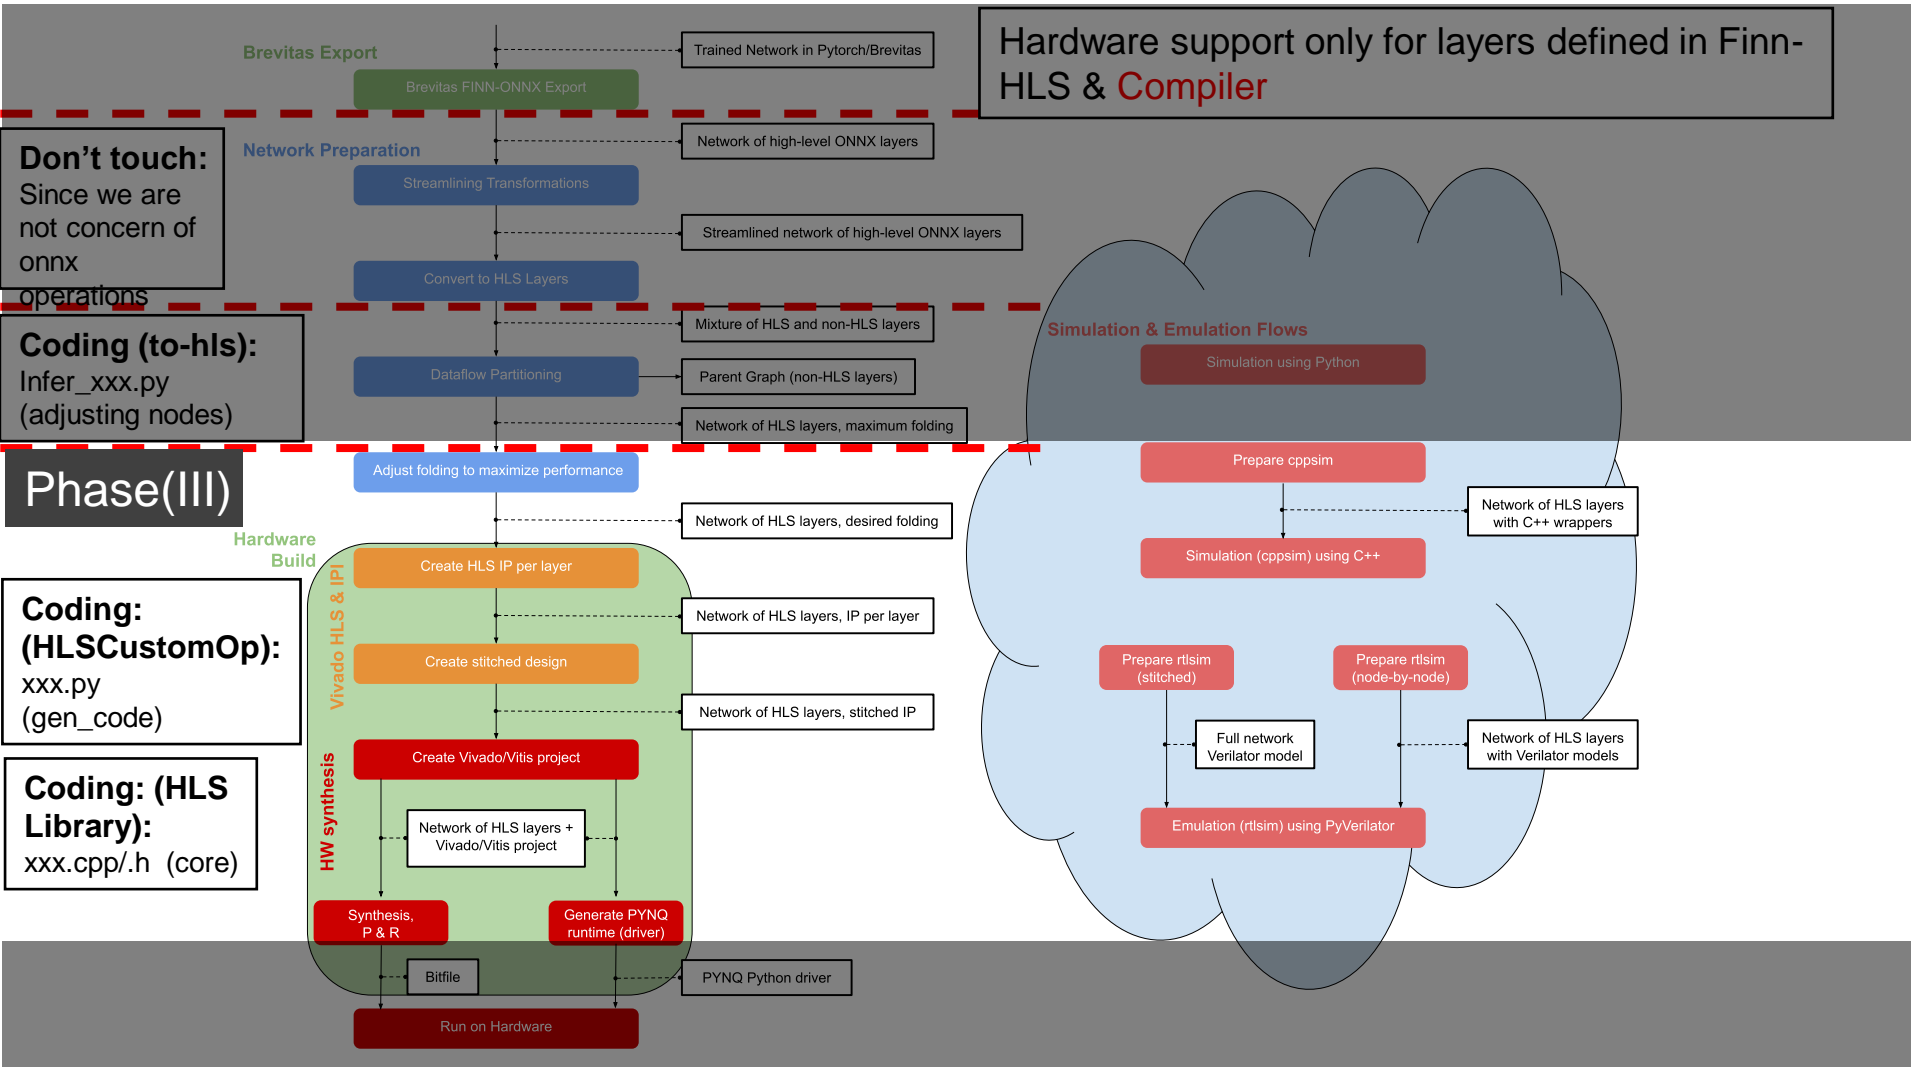


NODE PROPERTIES	
type	StreamingFCLayer_Batch
module	finn.custom_op.fpgadataflow
ATTRIBUTES	
accDataType	UINT10
ActVal	0
backend	fpgadataflow
binaryXnorMode	1
inFIFOdepth	16
inputDataType	BINARY
mem_mode	decoupled
MH	64
MW	784
noActivation	0
numInputVectors	1
outFIFOdepth	64
outputDataType	BINARY
PE	16
ram_style	block
SIMD	49
weightDataType	BINARY
INPUTS	

Overview

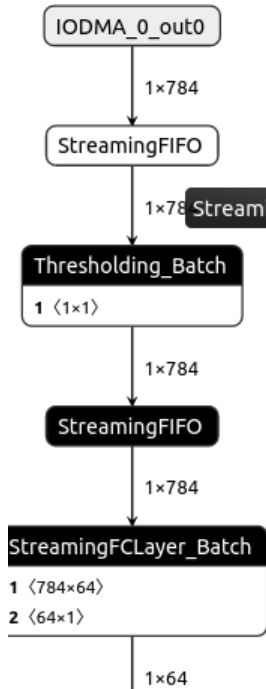
- End-to-End Compiling
- Phase (I): Brevitas export
- Phase (II):
 - Network preparation
 - Conversion to HLS layers
- **Phase (III): Hardware Build**
- Phase (IV): PYNQ deployment

Phase (III): Hardware Build



Phase (III): Hardware Build

- Automatically generate, link all the components and generate hardware bitstream
 - Zynq, Alveo -> ZynqBuild(), VitisBuild()
- ZynqBuild:
 - Preprocessing (Adding IODMA, DataWidthConverter, FIFO)
 1. PrepareIP
 2. HLSSynthIP
 3. CreateStitchedIP
 4. MakeZynqProject



NODE PROPERTIES

type	StreamingFIFO
module	finn.custom_op.fpgadataflow
name	StreamingDataflowPartition_1_StreamingFIFO_0

ATTRIBUTES

backend	fpgadataflow
code_gen_dir_ip...	/home/yuoto/multimedialC/FINN/practice/code/build_docker
code_gen_ipgen...	/code_gen_ipgen_StreamingDataflowPartition_1_StreamingFIFO_0_lr
dataType	UINT8
depth	32
folded_shape	1, 16, 49
ipgen_path	/home/yuoto/multimedialC/FINN/practice/code/bu /code_gen_ipgen_StreamingDataflowPartition_1_S /project_StreamingDataflowPartition_1_Streaming /verilog
ip_path	/home/yuoto/multimedialC/FINN/practice/code/bu /code_gen_ipgen_StreamingDataflowPartition_1_S /project_StreamingDataflowPartition_1_Streaming /verilog

code_gen_ipgen...IFO_0_lmbg7swt project_Streami...treamingFIFO_0 sol1 impl verilog

Name

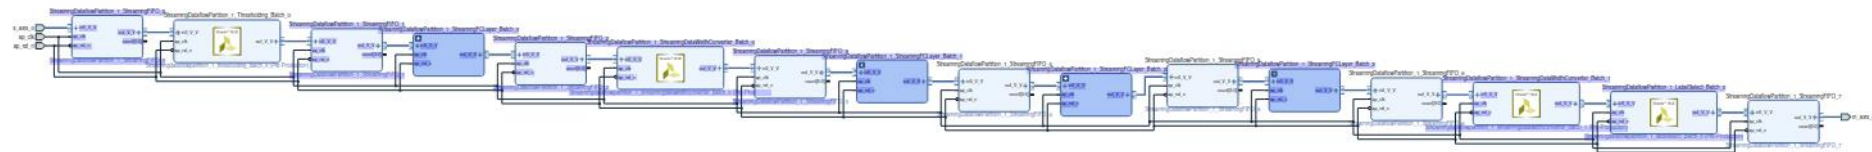
- xgui
- .hbs
- .Xil
- component.xml
- make_ip.sh
- package_ip.tcl
- Q_srl.v
- StreamingDataflowPartition_1_StreamingFIFO_0.v**
- StreamingDataflowPartition_1_StreamingFIFO_0.zip
- vivado.jou
- vivado.log

4. MakeZYNQProject(2/2)

- Vivado project was automatically built and stitched

```
model = ModelWrapper(postsynth_layers)
model.model.metadata_props
```

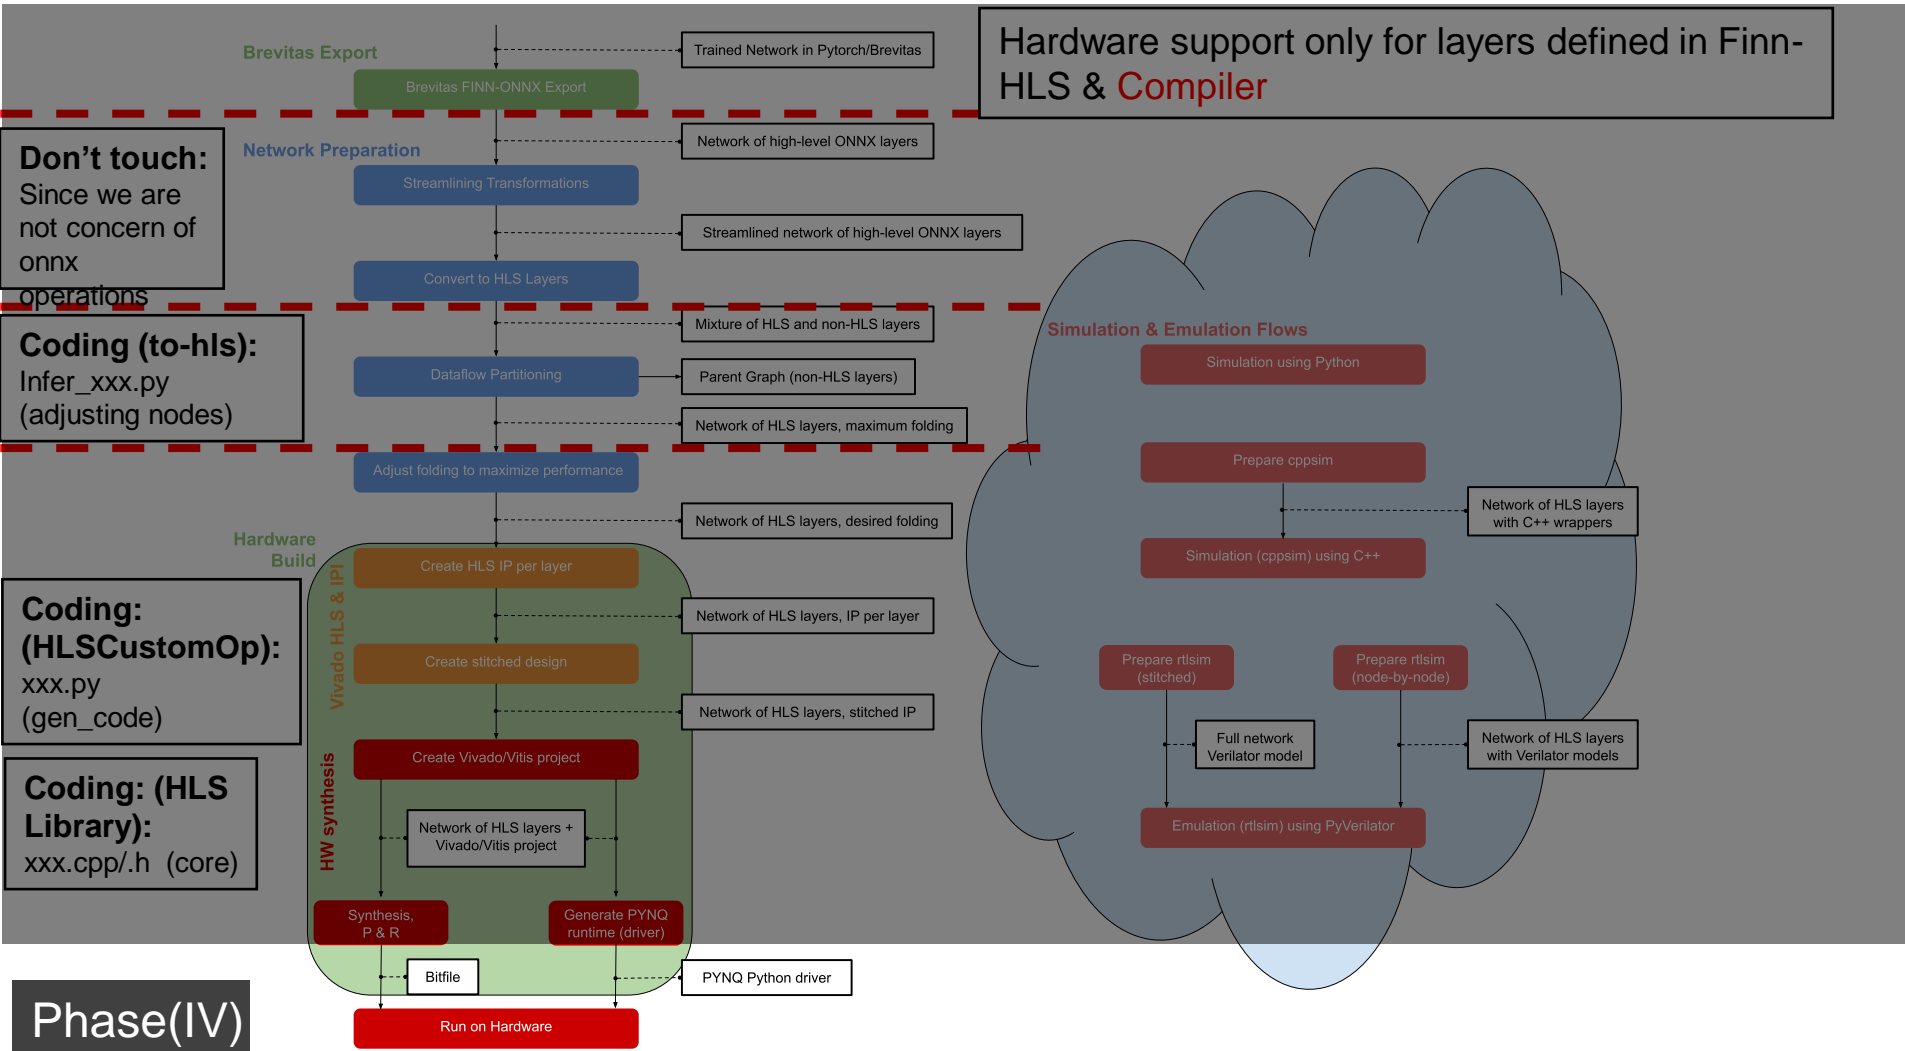
```
[key: "floorplan_json"
value: "/home/yuoto/multimediaIC/FINN/practice/code/build_docker/vitis_floorplan_zxz9em07/floorplan.json"
, key: "vivado_stitch_proj"
value: "/home/yuoto/multimediaIC/FINN/practice/code/build_docker/vivado_stitch_proj_ol2h0n49"
, key: "clk_ns"
value: "10"
, key: "wrapper_filename"
value: "/home/yuoto/multimediaIC/FINN/practice/code/build_docker/vivado_stitch_proj_ol2h0n49/finn_vivado_stitch_pr
oj.srcs/sources_1/bd/StreamingDataflowPartition_1/hdl/StreamingDataflowPartition_1_wrapper.v"
, key: "vivado_stitch_vlnv"
value: "xilinx_finn:finn:StreamingDataflowPartition_1:1.0"
, key: "vivado_stitch_ifnames"
value: "{\"clk\": [\"ap_clk\"], \"rst\": [\"ap_rst_n\"], \"s_axis\": [[\"s_axis_0\", 392]], \"m_axis\": [[\"m_axis
_0\", 8]], \"aximm\": [], \"axilite\": []}"
, key: "platform"
value: "zynq-iodma"
]
```



Overview

- End-to-End Compiling
- Phase (I): Brevitas export
- Phase (II):
 - Network preparation
 - Conversion to HLS layers
- Phase (III): Hardware Build
- Phase (IV): PYNQ deployment

Phase (IV): PYNQ deployment



Phase (IV): PYNQ deployment

- Deployment and Remote Execution
- Validating the Accuracy on a PYNQ Board
- Throughput Test on PYNQ Board