



Bridge of Life Education

FINN HLS

Lecturer: Ke-Han Li, Hua-Yang Weng

Date: 2022/4/27

[FPGA'17: FINN: A Framework for Fast, Scalable Binarized Neural Network Inference]
(<https://arxiv.org/abs/1612.07119>)

Outline

- Matrix-Vector Accumulation Unit
- Pooling Layer
- Convolution Layer
- FIFO

Outline

- Matrix-Vector Accumulation Unit
- Pooling Layer
- Convolution Layer
- FIFO


Matrix-Vector Accumulation Unit

- The function performs the multiplication between a weight matrix and the input activation vector.

$$\begin{bmatrix} 1 & 2 & 3 & 4 \\ 2 & 3 & 4 & 5 \\ 3 & 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix} = \begin{bmatrix} 30 \\ 40 \\ 50 \end{bmatrix}$$

Block Diagram

- According to the unroll factor(PE 、SIMD) set in python code, the compiler will generate corresponding amount of hardware.



I/O interface

- hls::stream &in
- hls::stream &out
- TW const &weights
- TA const &activation

Constant variables

- unsigned const $NF = \text{MatrixH} / \text{PE};$
- unsigned const $SF = \text{MatrixW} / \text{SIMD};$
- unsigned const $\text{TOTAL_FOLD} = NF * SF;$

HLS Code (1/)

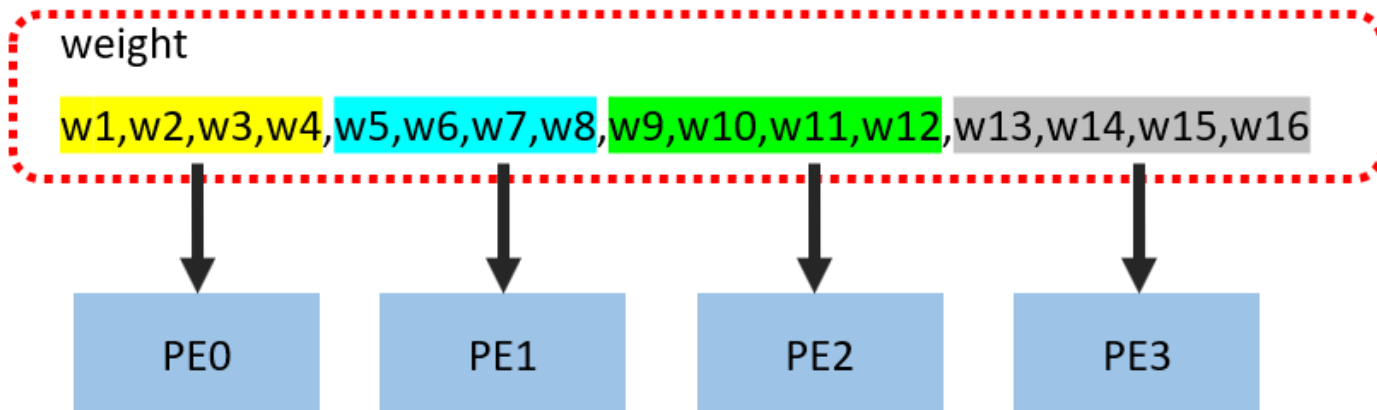
- MVAU reads the input activations from input port, and store them into internal buffer

```
248     for(unsigned i = 0; i < reps * TOTAL_FOLD; i++) {
249     #pragma HLS PIPELINE II=1
250         TI inElem;
251
252         if(nf == 0) {
253             // read input from stream
254             inElem = in.read();
255             // store in appropriate buffer for reuse
256             inputBuf[sf] = inElem;
257         }
258         else {
259             // reuse buffered input
260             inElem = inputBuf[sf];
261         }
```


HLS Code (2/)

- MVAU reads the weights and packed them.

```
263     // read from the parameter stream
264     W_packed = weight.read();
265     for (unsigned pe = 0; pe < PE; pe++) {
266 #pragma HLS UNROLL
267         w.m_weights[pe] = W_packed((pe+1)*SIMD*TW::width-1,pe*SIMD*TW::width);
268     }
```



HLS Code (3/)

- Set the accumulation register value to zero.

```
270      // Threshold Initialisation
271      if(sf == 0) {
272          for(unsigned pe = 0; pe < PE; pe++) {
273              #pragma HLS UNROLL
274                  accu[0][pe] = activation.init(nf, pe);
275              }
276          }
```

HLS Code (4/)

- MVAU uses mac unit to complete matrix vector calculation, and the number of mac unit depends on PE value.

```
278     // compute matrix-vector product for each processing element
279     for(unsigned pe = 0; pe < PE; pe++) {
280 #pragma HLS UNROLL
281         auto const act = TSrcI()(inElem, 0);
282         auto const wgt = TWeightI()(w[pe]);
283         //auto const wgt = w[pe];
284         accu[0][pe] = mac<SIMD>(accu[0][pe], wgt, act, r, 0);
285     }
```


HLS Code (5/)

- Feed the results into activation function.
- Write the activation function output into output buffer.


```
288     ++tile;
289     if(++sf == SF) {
290         // produce output and clear accumulators
291         auto outElem = TDstI().template operator()<TO>();
292         for (unsigned pe = 0; pe < PE; pe++) {
293 #pragma HLS UNROLL
294             outElem(pe,0,1) = activation.activate(nf, pe, accu[0][pe]);
295         }
296
297         out.write(outElem);
```

Example

Initial



x



=



Width:8

High:4

PE:2


SIMD:4

NF:2

SF:2

TOTAL_FOLD:4


Fully Connected Layer



Weights

Input

Output



Input

Weights

Output

StreamingFCLayer_Batch_3

```

5 // includes for network parameters
6 #include "weights.hpp"
7 #include "activations.hpp"
8 #include "mvau.hpp"
9
10 // defines for network parameters
11 #define MW1 64
12 #define MH1 10
13
14 #define SIMD1 8
15 #define PE1 10
16 #define WMEM1 8
17
18 #define TMEM1 0
19 #define numReps 1
20 #define WP1 1
21
22
23 void StreamingDataflowPartition_1_StreamingFCLayer_Batch_3(
24     hls::stream<ap_uint<8>> &in0,
25     hls::stream<ap_uint<80>> &weights,
26     hls::stream<ap_uint<80>> &out
27 )
28 {
29     #pragma HLS INTERFACE axis port=in0
30     #pragma HLS INTERFACE axis port=out
31     #pragma HLS stream depth=64 variable=in0
32     #pragma HLS stream depth=10 variable=out
33     #pragma HLS INTERFACE ap_ctrl_none port=return
34     #pragma HLS INTERFACE axis port=weights
35     #pragma HLS stream depth=8 variable=weights
36     Matrix_Vector_Activate_Stream_Batch<MW1, MH1, SIMD1, PE1, Recast<XnorMul>, Slice<ap_uint<8>>, Identity, ap_uint<1>> >
37     (in0, out, weights, PassThroughActivation<ap_uint<8>>(), numReps, ap_resource_lut());
38 }

```


I/O interface

- hls::stream<ap_uint<8>> &in0
- hls::stream<ap_uint<80>> &weights
- hls::stream<ap_uint<80>> &out

```
22
23 void StreamingDataflowPartition_1 StreamingECLayer_Batch_3(
24     hls::stream<ap_uint<8>> &in0,
25     hls::stream<ap_uint<80>> &weights,
26     hls::stream<ap_uint<80>> &out
27 )
28 {
29     #pragma HLS INTERFACE axis port=in0
30     #pragma HLS INTERFACE axis port=out
31     #pragma HLS stream depth=64 variable=in0
32     #pragma HLS stream depth=10 variable=out
33     #pragma HLS INTERFACE ap_ctrl_none port=return
34     #pragma HLS INTERFACE axis port=weights
35     #pragma HLS stream depth=8 variable=weights
36     Matrix_Vector_Activate_Stream_Batch<MW1, MH1, SIMD1, PE1, Recast<XnorMul>, Slice<ap_uint<8>>, Identity, ap_uint<1> >
37         (in0, out, weights, PassThroughActivation<ap_uint<8>>(), numReps, ap_resource_lut());
38 }
```

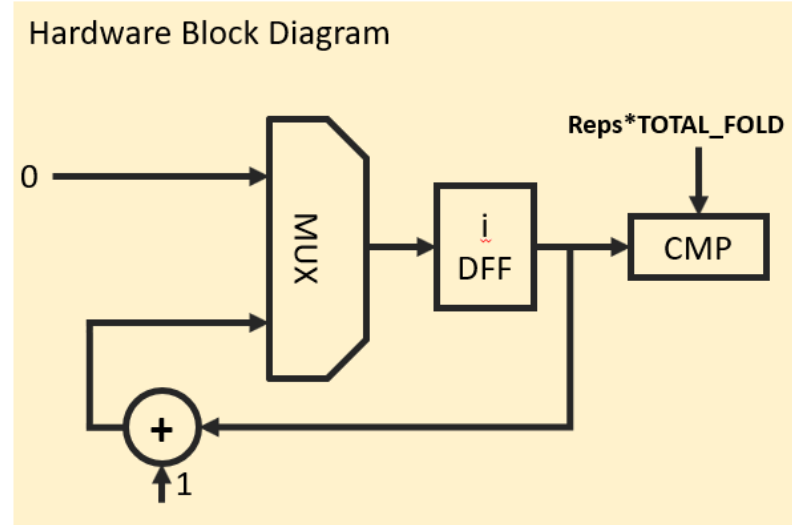
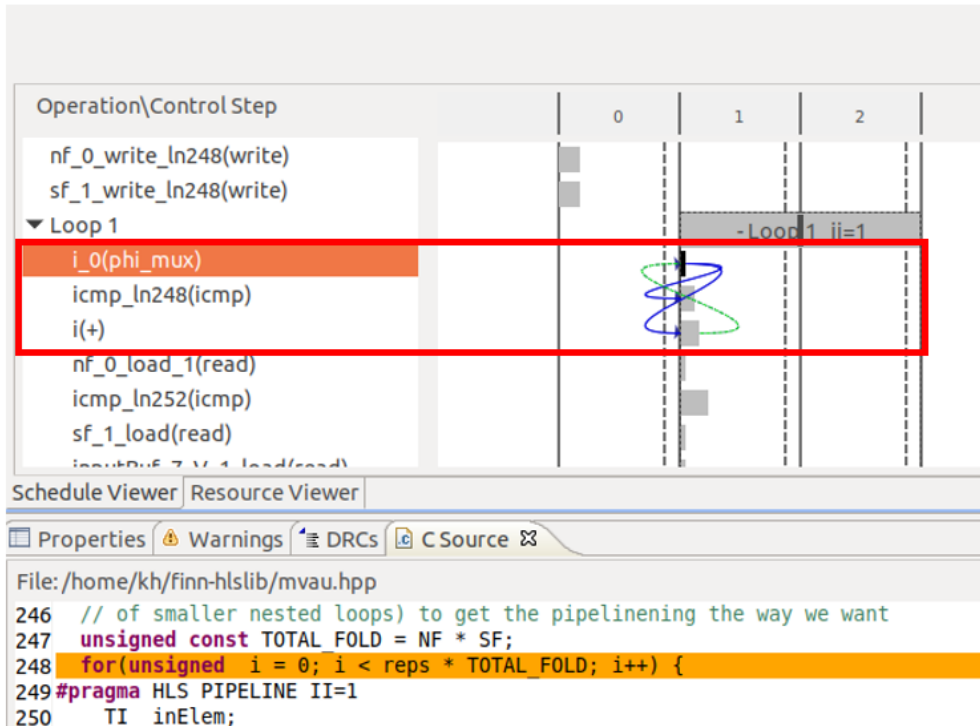

Latency Report

- The total fold is $(MH/PE) * (MW/SIMD)$.
- $(10/10)*(64/8)=8$.

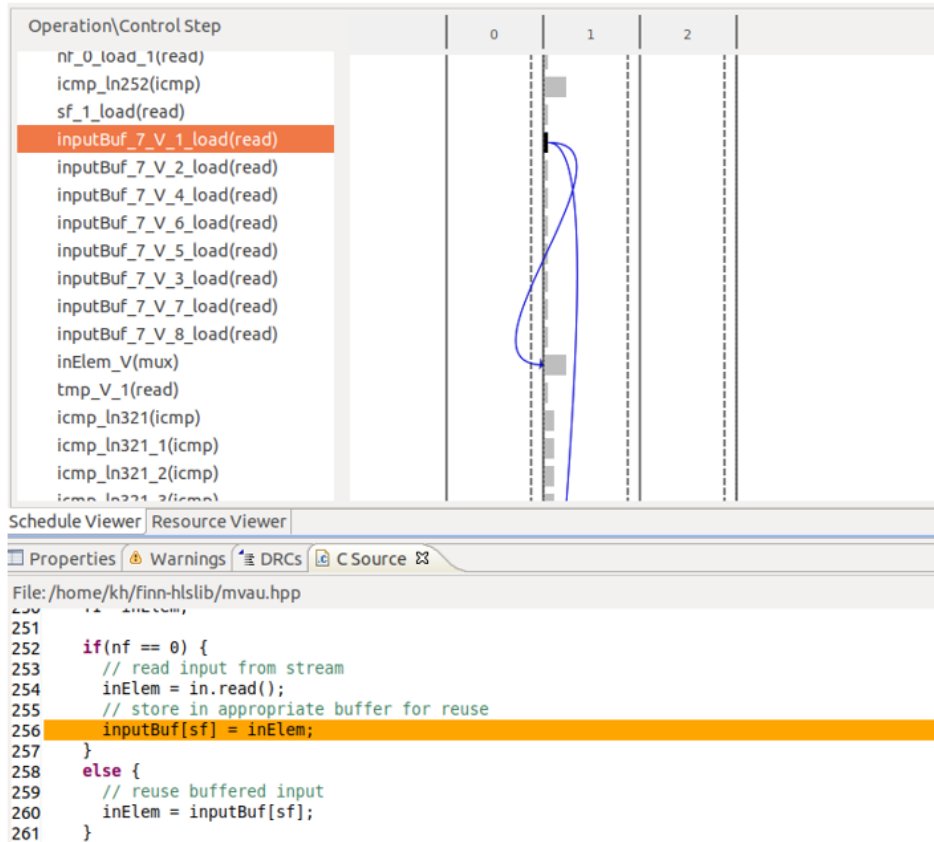


	Pipelined	Latency	Iteration Latency
Matrix_Vector_Activa	-	10	-
Loop 1	yes	8	2


Schedule Viewer Trace(1/)






Schedule Viewer Trace(2/)



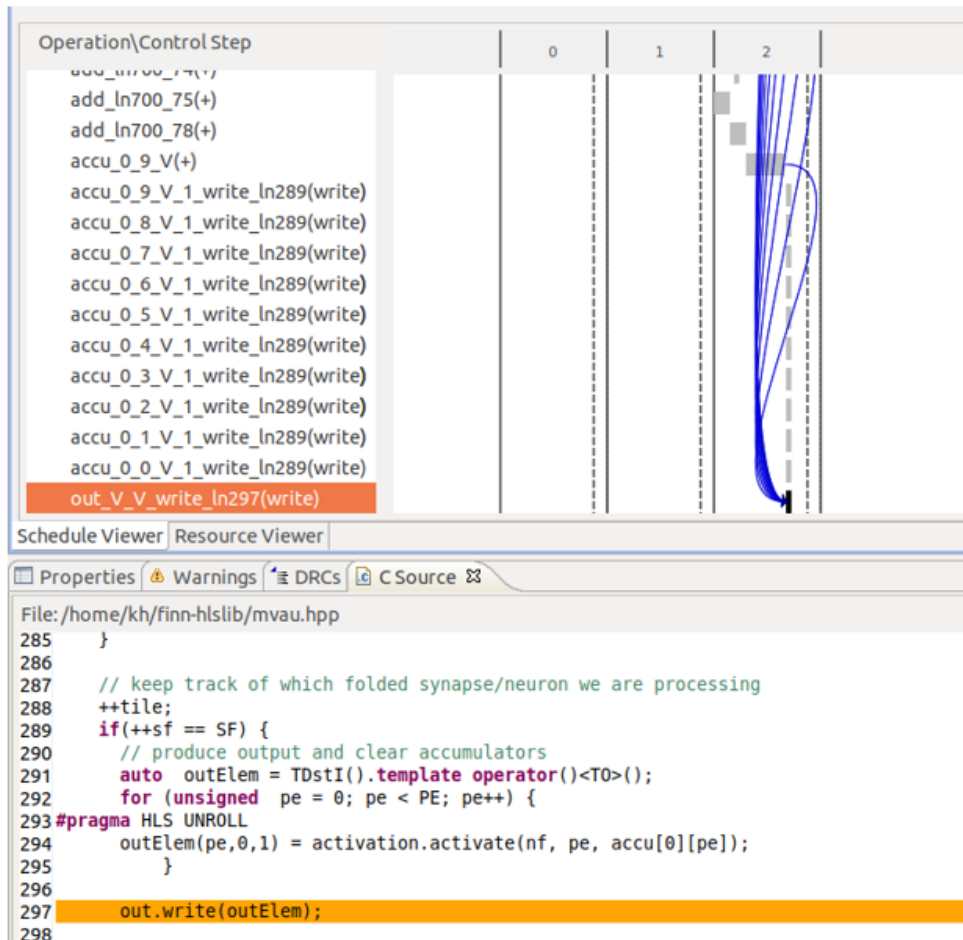
Hardware Block Diagram




Schedule Viewer Trace(3/)




Schedule Viewer Trace(4/)



PE0




PE_x



Outline



- Matrix-Vector Accumulation Unit
- Pooling Layer
- Convolution Layer
- FIFO

Max Pooling Layer



Binary Max Pooling Layer

We can use OR Gate to implement this operation.




HLS Code

```

67 void StreamingMaxPool(stream<ap_uint<NumChannels> > & in,
68     stream<ap_uint<NumChannels> > & out) {
69     CASSERT_DATAFLOW(ImgDim % PoolDim == 0);
70     // need buffer space for a single maxpooled row of the image
71     ap_uint<NumChannels> buf[ImgDim / PoolDim];
72     for(unsigned int i = 0; i < ImgDim / PoolDim; i++) {
73         #pragma HLS UNROLL
74         buf[i] = 0;
75     }
76
77     for (unsigned int yp = 0; yp < ImgDim / PoolDim; yp++) {
78         for (unsigned int ky = 0; ky < PoolDim; ky++) {
79             for (unsigned int xp = 0; xp < ImgDim / PoolDim; xp++) {
80                 #pragma HLS PIPELINE II=1
81                 ap_uint<NumChannels> acc = 0;
82                 for (unsigned int kx = 0; kx < PoolDim; kx++) {
83                     acc = acc | in.read();
84                 }
85                 // pool with old value in row buffer
86                 buf[xp] |= acc;
87             }
88         }
89         for (unsigned int outpix = 0; outpix < ImgDim / PoolDim; outpix++) {
90             #pragma HLS PIPELINE II=1
91             out.write(buf[outpix]);
92             // get buffer ready for next use
93             buf[outpix] = 0;
94         }
95     }
96 }


```

Example




Input Feature map: 8x8
Max pooling kernel: 2x2

Line buffer



output FIFO




Outline

- Matrix-Vector Accumulation Unit
- Pooling Layer
- Convolution Layer
- FIFO

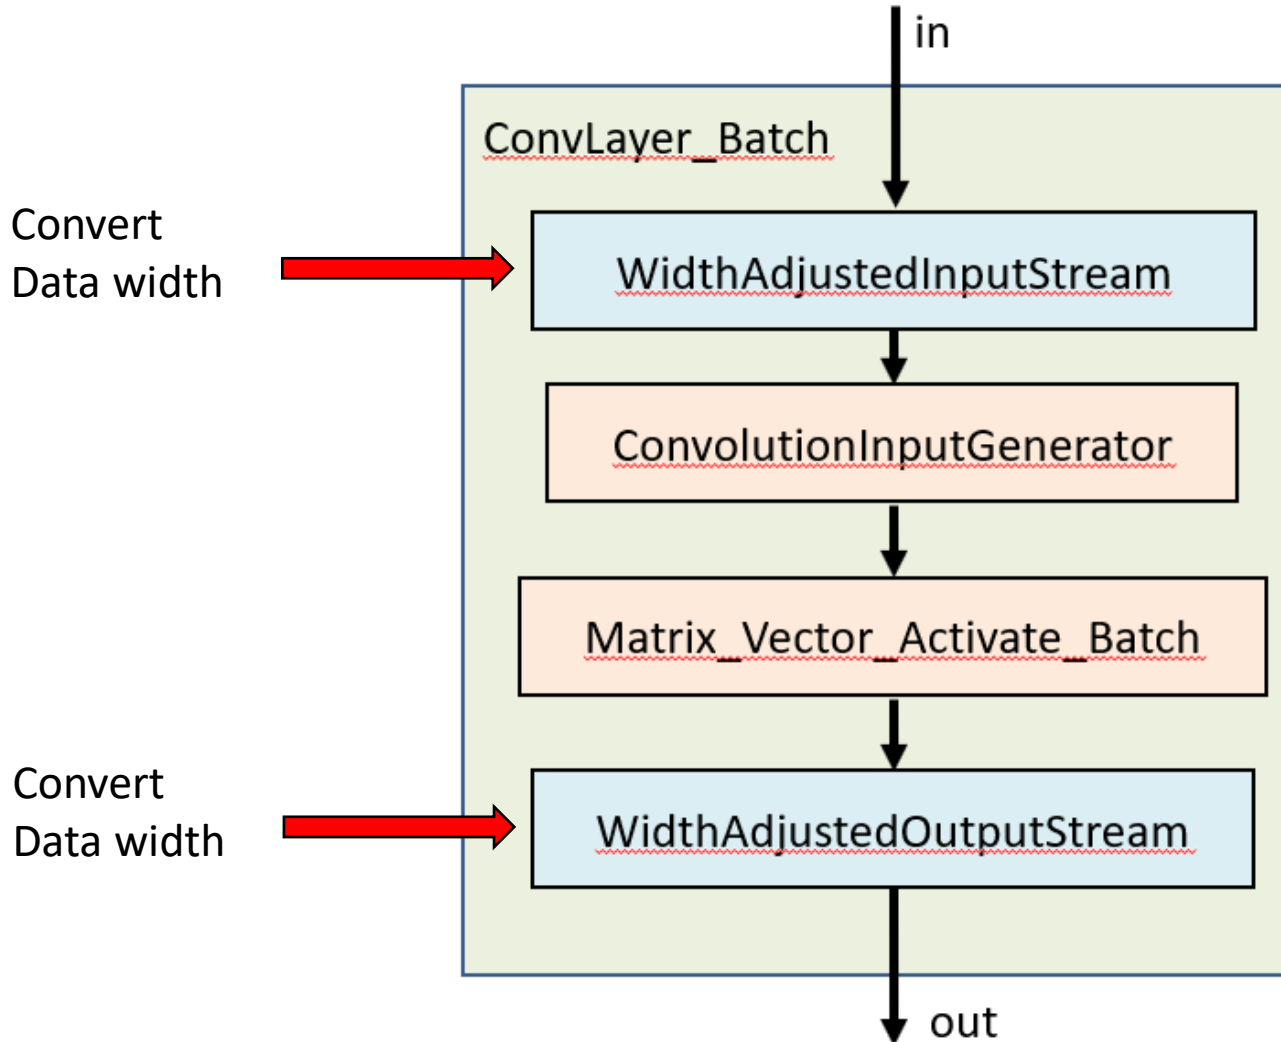
Convolution Layer

- Mainly consist of 2 core units
 - Convolution Input Generator



- Matrix Vector Activation Batch

Convolution Layer: Block Diagram



ConvLayer_Batch()


ConvLayer_batch():
ConvolutionInputGenerator (SWU)
Matrix_Vector_Activate_Batch(mvau)

```
template<
    unsigned int ConvKernelDim,
    unsigned int IFMChannels,
    unsigned int IFMDim,
    unsigned int OFMChannels,
    unsigned int OFMDim,

    unsigned int SIMD,          // number of SIMD lanes
    unsigned int PE,           // number of PEs

    typename TSrcI = Identity,  // redefine I/O interpretation as needed for input activations
    typename TDstI = Identity,  // redefine I/O interpretation as needed for output activations
    typename TWeightI = Identity, // redefine I/O interpretation as needed for weights

    int InStreamW, int OutStreamW, // safely deducible (stream width must be int though!)
    typename TW,   typename TA,   typename R
>
void ConvLayer_Batch(hls::stream<ap_uint<InStreamW>> &in,
                    hls::stream<ap_uint<OutStreamW>> &out,
                    TW const      &weights,
                    TA const      &activation,
                    unsigned const reps,
                    R const &r) {
    #pragma HLS INLINE
    unsigned const MatrixW = ConvKernelDim * ConvKernelDim * IFMChannels;
    unsigned const MatrixH = OFMChannels;
    unsigned const InpPerImage = IFMDim*IFMDim*IFMChannels*TSrcI::width/InStreamW;
    WidthAdjustedInputStream <InStreamW, SIMD*TSrcI::width, InpPerImage> wa_in (in, reps);
    WidthAdjustedOutputStream <PE*TDstI::width, OutStreamW, OFMDim * OFMDim * (OFMChannels / PE)> mvOut (out, reps);
    hls::stream<ap_uint<SIMD*TSrcI::width> > convInp("StreamingConvLayer_Batch.convInp");
    ConvolutionInputGenerator<ConvKernelDim, IFMChannels, TSrcI::width, IFMDim,
                             OFMDim, SIMD,1>(wa_in, convInp, reps, ap_resource_dflt());
    Matrix_Vector_Activate_Batch<MatrixW, MatrixH, SIMD, PE, 1, TSrcI, TDstI, TWeightI>
    (static_cast<hls::stream<ap_uint<SIMD*TSrcI::width>>&>(convInp),
     static_cast<hls::stream<ap_uint<PE*TDstI::width>>&> (mvOut),
     weights, activation, reps* OFMDim * OFMDim, r);
}
```



NODE PROPERTIES

type StreamingDataflowPartition

module finn.custom_op.fpgadataflow

name StreamingDataflowPartition_1

ATTRIBUTES

mem_port +

model /home/yuoto/multimedialC/FINN/finn_new_vitis/host_build/dataflow_partition2_wbswq3d3/df_model.onnx +

partition_id 2 +


slr -1 +

INPUTS

0 name: StreamingDataflowPartition_0_out0 +

OUTPUTS

0 name: StreamingDataflowPartition_1_out0 +



NODE PROPERTIES	
ATTRIBUTES	
backend	fpgadataflow
code_gen_dir_ip ...	/home/yuoto/multimedialC/FINN/finn_new_vitis/host_build /code_gen_ipgen_StreamingDataflowPartition_1_ConvolutionInput
ConvKernelDim	3, 3
depthwise	0
Dilation	1, 1
IFMChannels	3
IFMDim	32, 32
inputDataType	INT8
StreamingDataflowPartition_1_ConvolutionInputGenerator_0	
ip_path	/home/yuoto/multimedialC/FINN/finn_new_vitis/host_build /code_gen_ipgen_StreamingDataflowPartition_1_ConvolutionInput /project_StreamingDataflowPartition_1_ConvolutionInputGeneratc
ip_vlnv	xilinx.com:hls:StreamingDataflowPartition_1_ConvolutionInputGen
OFMDim	30, 30
outFIFODepth	128
outputDataType	INT8
partition_id	2

- ConvKernelDim: kernel size [3, 3]
- Dilation: [1, 1]
- Stride:[1, 1]
- IFMChannels(Input Feature Map): 3 (Int8 RGB)
- IFMDIM: [32, 32]
- SIMD: 3

ConvolutionInputGenerator



Stride = 1
IFMDim = 32
OFMDim = 30
multiplying_factor = 1
number_blocks = 4
cycle_write_block = 270
cycle_read_block = 32

counter_internal_block = 0
current_line = 0
current_block_read = 0
current_line_in_block = 0
count_simd = 0
current_block_write = 3
read_block = 3
kx = 0, ky = 0
ofm_x = 0, ofm_y = 0

ConvolutionInputGenerator(1/4)

```

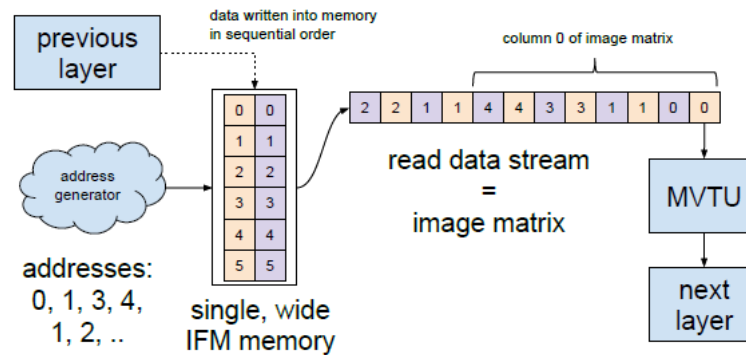
void ConvolutionInputGenerator(
    stream<ap_uint<SIMD*Input_precision> > & in,
    stream<ap_uint<SIMD*Input_precision> > & out,
    const unsigned int numReps,
    R const &r) {
    CASSERT_DATAFLOW(IFMChannels % SIMD == 0);
    CASSERT_DATAFLOW(ConvKernelDim % Stride == 0);
    const unsigned int multiplying_factor = IFMChannels/SIMD;
    const unsigned int number_blocks = ConvKernelDim/Stride + 1 ;
    ap_uint<SIMD*Input_precision> inputBuf[number_blocks][Stride * IFMDim * multiplying_factor];
#pragma HLS ARRAY_PARTITION variable=inputBuf complete dim=1
    memory_resource(inputBuf, r);
    const unsigned int cycles_write_block = (OFMDim * ConvKernelDim * ConvKernelDim * multiplying_factor);
    const unsigned int cycles_read_block = Stride * IFMDim * multiplying_factor;
    const unsigned int max_cycles = MAX(cycles_write_block, cycles_read_block);
    const unsigned int baseIter = IFMDim * ConvKernelDim * multiplying_factor// Initial buffer
        + OFMDim * MAX(cycles_write_block, cycles_read_block);
    unsigned int counter_internal_block = 0;
    unsigned int current_block_write = 0;
    unsigned int next_block_write = 0;
    unsigned int current_line = 0;
    unsigned int read_block = 0;
    unsigned int inp = 0, ofm_y = 0, ofm_x = 0, k_y = 0, k_x = 0, count_simd = 0;
    }
    template<unsigned int ConvKernelDim,
            unsigned int IFMChannels,
            unsigned int Input_precision,
            unsigned int IFMDim,
            unsigned int OFMDim,
            unsigned int SIMD,
            unsigned int Stride,
            typename R>

```


Define vars

ConvolutionInputGenerator

- **Outer loop:** *BaseIteration* to finish one image matrix
 - (if) First initialize input buffer
 - (else) Then ... parallel do
 - write the output in the correct order
 - Read the next block
 - Update count



(b) SWU operation.



- Goal: pipelined with $II=1$

ConvolutionInputGenerator(1/4)

Base Iteration:

(if) First initialize input buffer
(else) Then parallel do 3 ifs

```
#pragma HLS reset variable=inp
for (unsigned int count_image = 0; count_image < numReps; count_image++) {
    for (unsigned int i = 0; i < baseIter; i++) {
#pragma HLS PIPELINE II=1
        if (inp < IFMDim * ConvKernelDim*multiplying_factor) {// Initial buffer of ConvKe
            ap_uint<SIMD*Input_precision> inElem;
            inElem = in.read();
            inputBuf[current_block_write][current_line] = inElem;
            current_line++;
            inp++;
            if (current_line == Stride * IFMDim * multiplying_factor ) {
                current_line = 0;
                current_block_write++;
                if (current_block_write == number_blocks) {
                    current_block_write=0;
                }
                read_block++;
                counter_internal_block = 0;
            }
        } else {
```

IFMDim*factor*stride

kernelDim
+1

ConvolutionInputGenerator(1/4)

```

} else {
    if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
        unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
        if (current_block_read >= number_blocks) {
            current_block_read -= number_blocks;
        }
        unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
        ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
        out.write(outElem);
        count_simd++;
    }
    if (count_simd == multiplying_factor) {
        count_simd=0;
        k_x++;
        if (k_x == ConvKernelDim) {
            k_x = 0;
            k_y++;
            if (k_y == ConvKernelDim) {
                k_y = 0;
                ofm_x ++;
                if (ofm_x == OFMDim) {
                    ofm_x = 0;
                    ofm_y++;
                    if (ofm_y == OFMDim) {
                        ofm_y = 0;
                        inp = 0;
                    }
                }
            }
        }
    }
}

```

1. if: write output
(within # of write cycles)

ConvolutionInputGenerator(4/4)

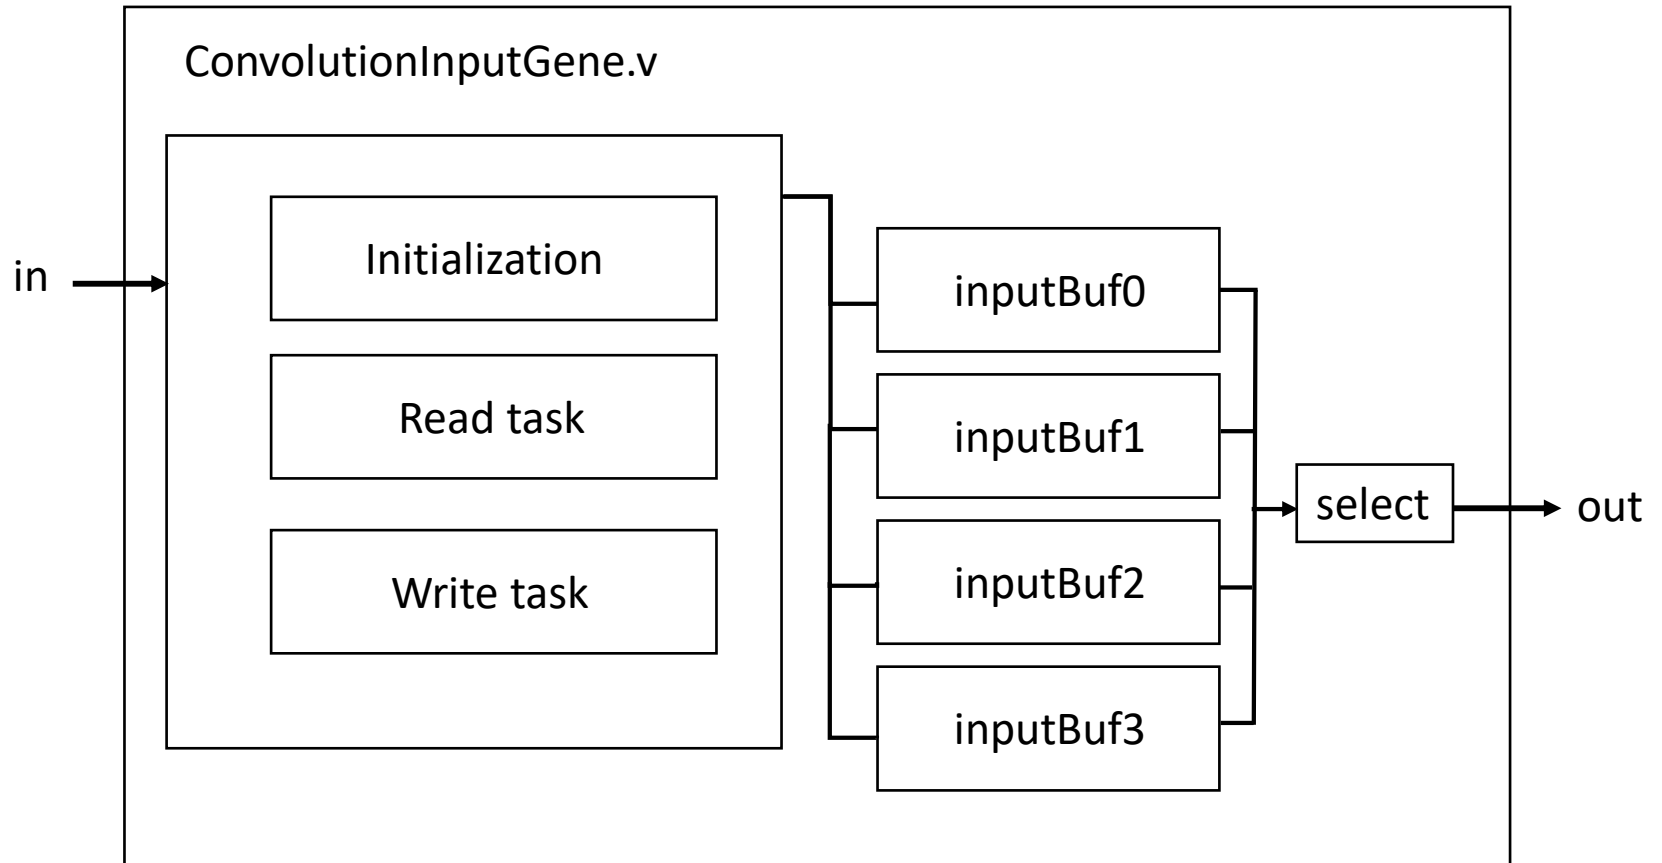
```


if ((counter_internal_block < cycles_read_block-1) && (read_block<IFMDim/Stride)) { //
    ap_uint<SIMD*Input_precision> inElem;
    inElem = in.read();
    inputBuf[current_block_write][current_line] = inElem;
#pragma AP dependence variable=inputBuf intra false
#pragma AP dependence variable=inputBuf inter false
    current_line++;
    if (current_line == Stride * IFMDim * multiplying_factor) {// We read the whole block
        // We filled up a block, let's not read until
        current_line = 0;
        read_block++;
        current_block_write++;
        if (current_block_write == number_blocks) {
            current_block_write=0;
        }
#pragma AP dependence variable=current_block_write intra false
    }
}
counter_internal_block++; // = (counter_internal_block +1) % max_cycles;
if (counter_internal_block == (max_cycles-1)) {
    counter_internal_block = 0;
}
}
} // End base_iter
read_block = 0;
} // End count_image
} // End generator

```

2. if: read input to buffer
(within # of read cycles and image)

3. if: initialize the counter when one R/W
pair is achieved. (max_cycle is reached)



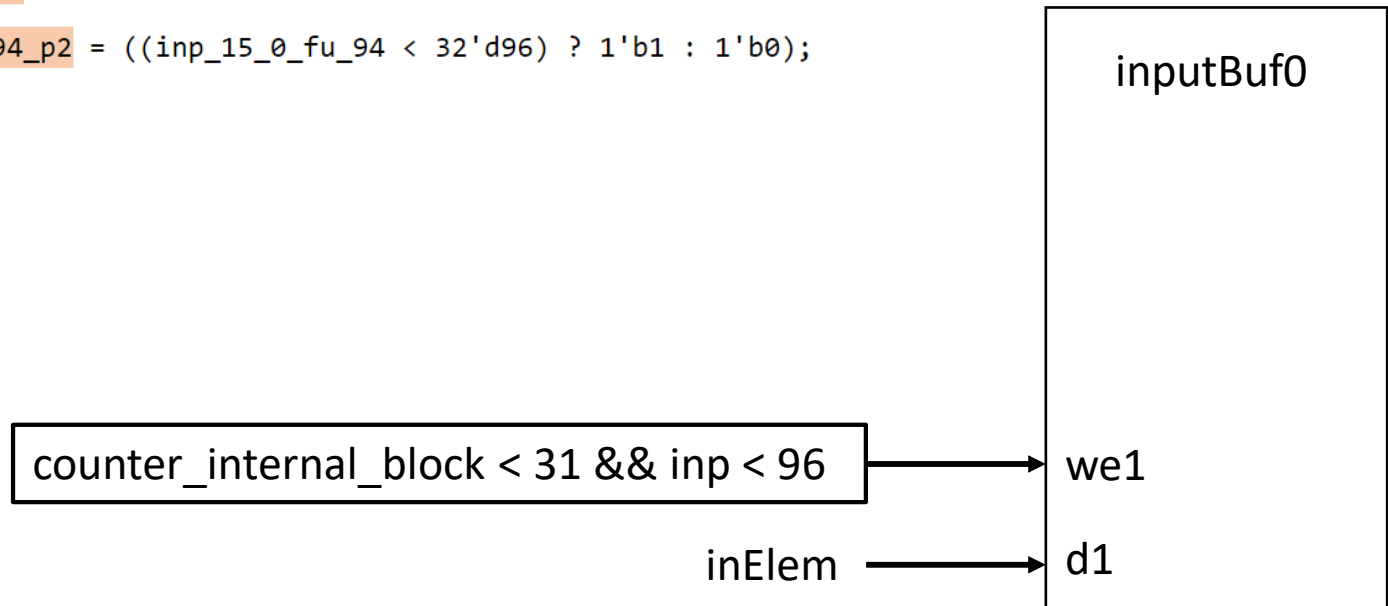



```
always @ (*) begin
    if (((1'b0 == ap_block_pp0_stage0_11001) & (1'd1 == and_ln244_fu_598_p2) & (icmp_ln198_fu_394_p2 == 1'd0) & (icmp_
        | inputBuf_0_V_we1 = 1'b1;
    end else begin
        | inputBuf_0_V_we1 = 1'b0;
    end
end
```


```
assign and_ln244_fu_598_p2 = (icmp_ln244_fu_576_p2 & icmp_ln244_1_fu_592_p2);
```

```
assign icmp_ln244_fu_576_p2 = ((counter_internal_blo_fu_118 < 32'd31) ? 1'b1 : 1'b0);
```


```
assign icmp_ln198_fu_394_p2 = ((inp_15_0_fu_94 < 32'd96) ? 1'b1 : 1'b0);
```




Case Study



Input IFM
Width:7
High:7
Channel:1
SIMD:1
Stride:1



Kernel:3x3



Output OFM
Width:5
High:5
Channel:1

$\text{multiplying_factor} = \text{IFMChannels} / \text{SIMD}$

(# of cycles to pick all data channel-wise)

Case Study: inputBuf

- `ap_uint<SIMD*Input_precision>`
 - `inputBuf[number_blocks][Stride * IFMDim * multiplying_factor]`
- `number_blocks = ConvKernelDim/Stride + 1 = 4`
- `Stride * IFMDim * multiplying_factor = 7`
- `ap_uint<SIMD*Input_precision> inputBuf[4][7]`

Case Study: baselter

- $\text{baselter} = \text{IFMDim} * \text{ConvKernelDim} * \text{multiplying_factor} // \text{Initial} + \text{OFMDim} * \text{MAX}(\text{cycles_write_block}, \text{cycles_read_block});$
- $\text{IFMDim} * \text{ConvKernelDim} * \text{multiplying_factor} = 21$
- $\text{OFMDim} * \text{MAX}(\text{cycles_write_block}, \text{cycles_read_block}) = 225$

i=0
inp=0

Input IFM


```

198 ✓ if (inp < IFMDim * ConvKernelDim*multiplying_factor) { // Initial buffer of ConvKernelDim lines
199   ap_uint<SIMD*Input_precision> inElem;
200   inElem = in.read();
201   inputBuf[current_block_write][current_line] = inElem;
202   current_line++;
203   inp++;
204 ✓   if (current_line == Stride * IFMDim * multiplying_factor ) {
205     current_line = 0;
206     current_block_write++;
207 ✓     if (current_block_write == number_blocks) {
208       current_block_write=0;
209     }
210     read_block++;
211     counter_internal_block = 0;
212   }
213 ✓ } else {

```

InputBuff

i=6
inp=6

Input IFM


```

198 ✓ if (inp < IFMDim * ConvKernelDim*multiplying_factor) { // Initial buffer of ConvKernelDim lines
199   ap_uint<SIMD*Input_precision> inElem;
200   inElem = in.read();
201   inputBuf[current_block_write][current_line] = inElem;
202   current_line++;
203   inp++;
204 ✓   if (current_line == Stride * IFMDim * multiplying_factor ) {
205     current_line = 0;
206     current_block_write++;
207 ✓     if (current_block_write == number_blocks) {
208       current_block_write=0;
209     }
210     read_block++;
211     counter_internal_block = 0;
212   }
213 ✓ } else {

```

InputBuff

$i=20$
 $inp=20$

Input IFM

$$7*3*1=21$$

```

198 ✓ if (inp < IFMDim * ConvKernelDim*multiplying_factor) { // Initial buffer of ConvKernelDim lines
199   ap_uint<SIMD*Input_precision> inElem;
200   inElem = in.read();
201   inputBuf[current_block_write][current_line] = inElem;
202   current_line++;
203   inp++;
204 ✓   if (current_line == Stride * IFMDim * multiplying_factor ) {
205     current_line = 0;
206     current_block_write++;
207 ✓     if (current_block_write == number_blocks) {
208       current_block_write=0;
209     }
210     read_block++;
211     counter_internal_block = 0;
212   }
213 ✓ } else {

```

InputBuff

i=21(do else)

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read = number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;

```

0

0

Kx=0
Ky=0
Ox=0
Oy=0

InputBuff



SWU output (image matrix)


```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read = number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y*Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;

```

SWU output (image matrix)

i=23

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read -= number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;

```

0

2

Kx=2
Ky=0
Ox=0
Oy=0

InputBuff



SWU output (image matrix)

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read = number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y*Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;

```

1


0

Kx=0

Ky=1

Ox=0

Oy=0



SWU output (image matrix)

i=29

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read -= number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;


```

2

2

Kx=2
Ky=2
Ox=0
Oy=0

InputBuff



SWU output (image matrix)

i=30

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read -= number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;


```

0

1

Kx=0
Ky=0
Ox=1
Oy=0

InputBuff



i=65

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read = number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;


```

2

6

Kx=2
Ky=2
Ox=4
Oy=0

InputBuff



i=66

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read = number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;

```

1

0

Kx=0
Ky=0
Ox=0
Oy=1

InputBuff

i=72

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read -= number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;

```

3

0

Kx=0
Ky=2
Ox=0
Oy=1

InputBuff

i=111

```

213 } else {
214     if (counter_internal_block < cycles_write_block-1) { // We are writing output, MMV IFMChan per cycle
215         unsigned int current_block_read = (current_block_write + 1 + k_y / Stride);
216         if (current_block_read >= number_blocks) {
217             current_block_read -= number_blocks;
218         }
219         unsigned int current_line_in_block = ((k_y%Stride) * IFMDim + ofm_x*Stride + k_x)*multiplying_factor + count_simd;
220         ap_uint<SIMD*Input_precision> outElem = inputBuf[current_block_read][(current_line_in_block)];
221         out.write(outElem);
222         count_simd++;
223         if (count_simd == multiplying_factor) {
224             count_simd=0;
225             k_x++;
226             if (k_x == ConvKernelDim) {
227                 k_x = 0;
228                 k_y++;
229                 if (k_y == ConvKernelDim) {
230                     k_y = 0;
231                     ofm_x ++;
232                     if (ofm_x == OFMDim) {
233                         ofm_x = 0;
234                         ofm_y++;
235                         if (ofm_y == OFMDim) {
236                             ofm_y = 0;
237                             inp = 0;

```

2

0

Kx=0
Ky=0
Ox=0
Oy=2

InputBuff

Outline

- Matrix-Vector Accumulation Unit
- Pooling Layer
- Convolution Layer
- FIFO

FIFO

- Critical in the whole system design.
 - If the depth of the FIFO is not enough
 - Streaming process might encounter a stall in response to the "full" or "empty" signal.
 - The design at both sides cannot either read or write data into this FIFO, resulting in a decrease of the throughput.
- *hls::stream* interface
 - *ap_fifo*
 - Tool will make a FIFO with a default depth of 2.

```

fc_layers = model.get_nodes_by_op_type("StreamingFCLayer_Batch")
# (PE, SIMD, in_fifo_depth, out_fifo_depth, ramstyle) for each layer
config = [
    (16, 49, 16, 64, "block"),
    (8, 8, 64, 64, "auto"),
    (8, 8, 64, 64, "auto"),
    (10, 8, 64, 10, "distributed"),
]
for fcl, (pe, simd, ififo, ofifo, ramstyle) in zip(fc_layers, config):
    fcl_inst = getCustomOp(fcl)
    fcl_inst.set_nodeattr("PE", pe)
    fcl_inst.set_nodeattr("SIMD", simd)
    fcl_inst.set_nodeattr("inFIFODepth", ififo)
    fcl_inst.set_nodeattr("outFIFODepth", ofifo)
    fcl_inst.set_nodeattr("ram_style", ramstyle)

# set parallelism for input quantizer to be same as first layer's SIMD
inp_qnt_node = model.get_nodes_by_op_type("Thresholding_Batch")[0]
inp_qnt = getCustomOp(inp_qnt_node)
inp_qnt.set_nodeattr("PE", 49)

```

How to determine FIFO depths?

- Depends on the IO protocol, bandwidth, and the throughput of the data producer/consumer.
 - All these possible factors result in the difficulty of analytical FIFO depth calculation.
 - Blindly increase the depth to a great extent, if the budget of the area is infinite
- Empirically estimate suitable FIFO depths.
 1. Set all FIFO depths to a large number.
 2. Prepare enough test patterns that can simulate the actual data processing flow
 3. Record the maximum data occupancy of this FIFO.
 4. After the simulation is done, change the size of the FIFO to this empirical depth parameter.
 5. Measure the resulting throughput drop of the system with this new configuration.
 6. If no throughput drop, a more aggressive FIFO depth reduction method may use.

FIFO Depth in FINN

- FINN Compiler (hardware build *ZynqBuild()* phase)
 - *InsertFIFO()* inserts FIFOs with the manually specified depth.
 - FINN also provides *InsertAndSetFIFODepths()*.
- *InsertAndSetFIFODepths()*
 - First set all FIFO depths to a number of 16K and sends random input image patterns.
 - PyVerilator is used for testing the maximum FIFO occupancy.

InsertAndSetFIFO Depths()

```
class InsertAndSetFIFO Depths(Transformation):
```

```
    """Insert appropriate-depth StreamingFIFOs through RTLSim that preserve
    throughput in the created accelerator.
```

```
    Constructor arguments:
```

- clk_ns : clock period (used for IP preparation)
- max_qsrl_depth : FIFOs deeper than this will use Vivado IP instead of Verilog FIFOs (Q_srl.v)
- max_depth : how deep the "max"-sized FIFOs initially inserted will be
- swg_exception : call CapConvolutionFIFO Depths to make convolution FIFOs smaller where appropriate
- vivado_ram_style : the StreamingFIFO.ram_style attribute to be used for large FIFOs implemented by Vivado

```
    Assumed input graph properties:
```

- all nodes are fpgadataflow nodes
- no FIFOs inserted,
- (inFIFO Depth/outFIFO Depth attrs will be ignored)

```
    Output:
```

- graph with appropriate-depth FIFOs inserted

```
    Background:
```

```
    Even with all FINN HLS fpgadatflow layers appropriately parallelized, it is
    necessary to insert FIFOs between them to prevent stalls due to bursty
    behavior. The sizes of those FIFOs are hard to predict analytically, so
    we do the following:
```

- insert very deep (default 16k deep) FIFOs between all fpgadataflow nodes
- create stitched design
- run through rtlsim with stream of multiple random input images (to fill pipeline)
- keep track of observed maximum occupancy for each FIFO during rtlsim
- when sim finished, update each FIFO depth to maximum observed occupancy and set inFIFO Depth/outFIFO Depth attrs to 0 on relevant nodes

```
    """
```

✂ **Note:** *ZynqBuild()* class does not default use this FIFO insertion method

- We have to manually replace it.

```
# change this
kernel_model = kernel_model.transform(InsertFIFO())

# to this
kernel_model = kernel_model.transform(InsertAndSetFIFO Depths(fpgapart=self.fpga_part,
                                                                clk_ns=self.period_ns))
```